

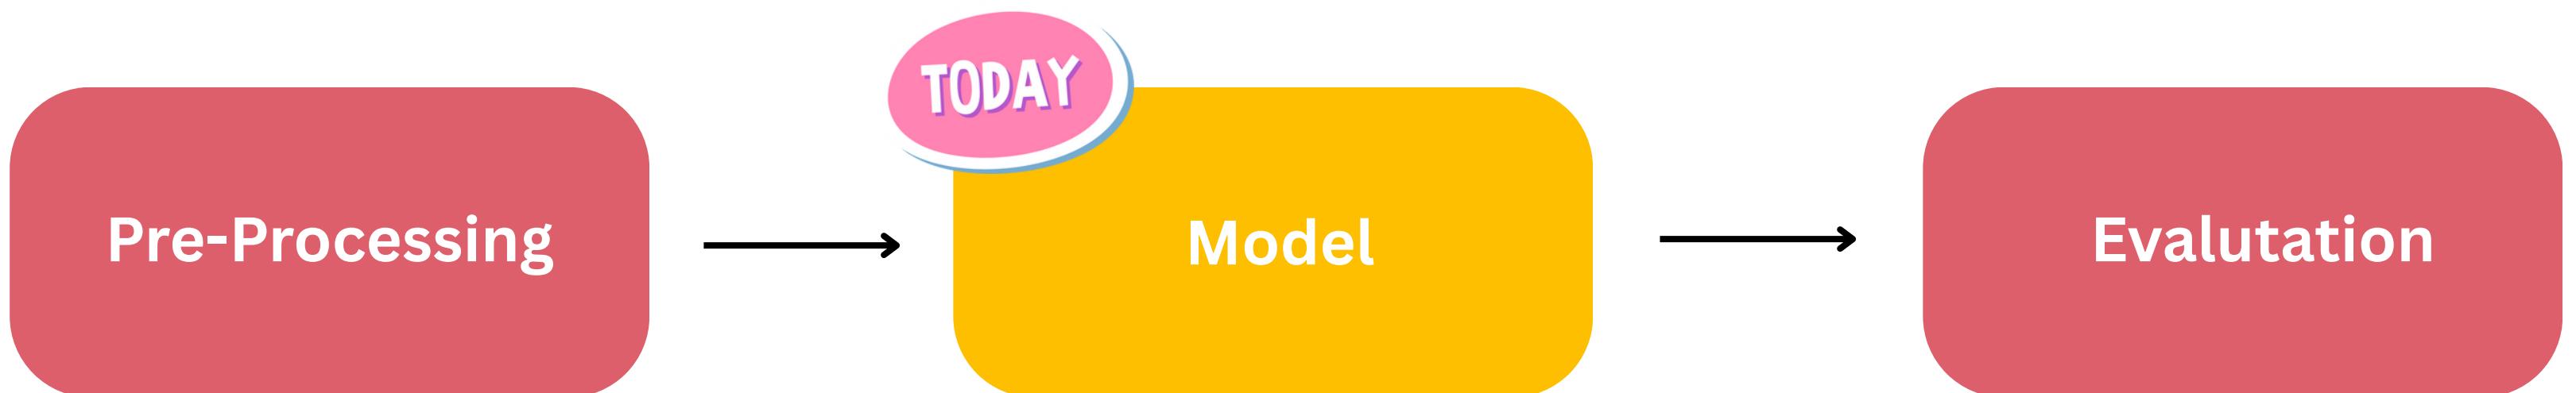


# Data Science and Machine Learning

For New Enthusiasts

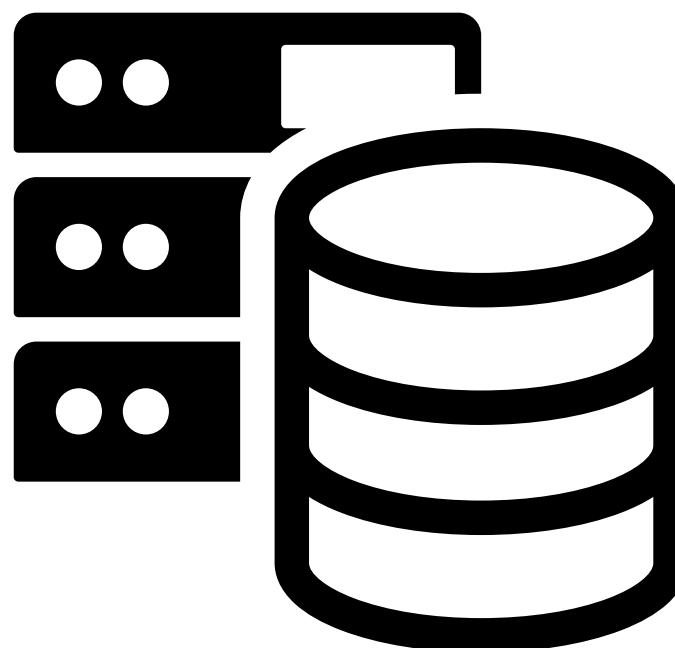
# REPO LINK





- **K-Nearest Neighbours**
- **Naive Bayes**
- **Decision Trees**
- **Artificial Neural Networks**

# What is Machine Learning?



**MODEL REPRESENTATION**

$f(x;\Theta)$ : Predictors  $\rightarrow$  Target

**Goal**



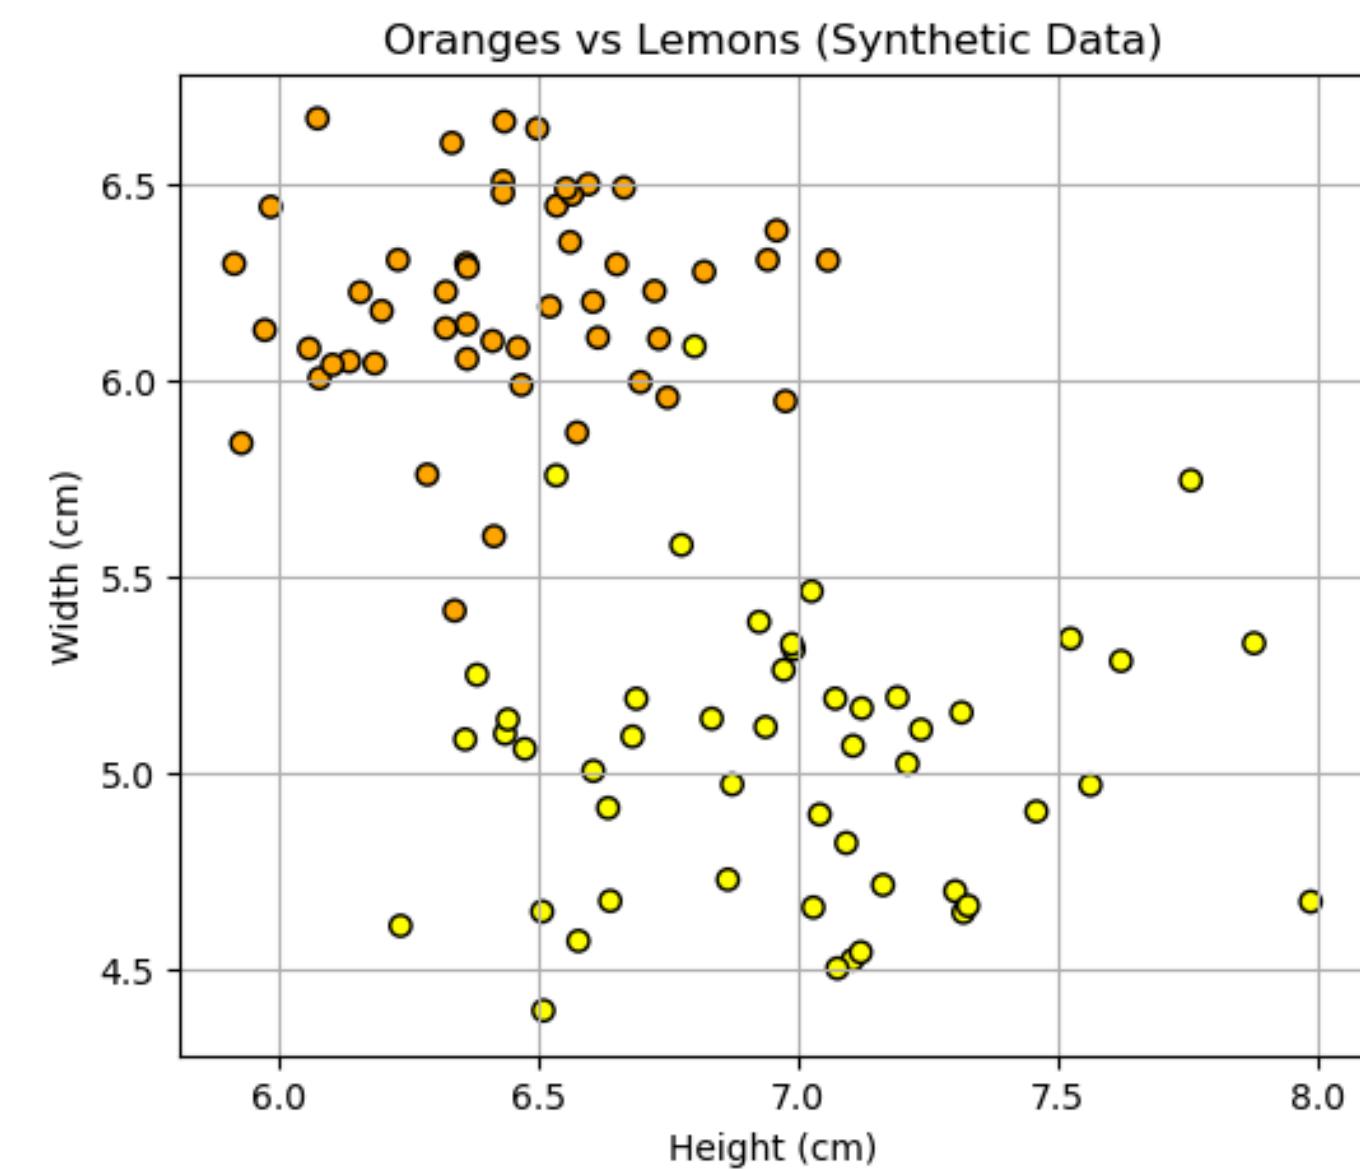
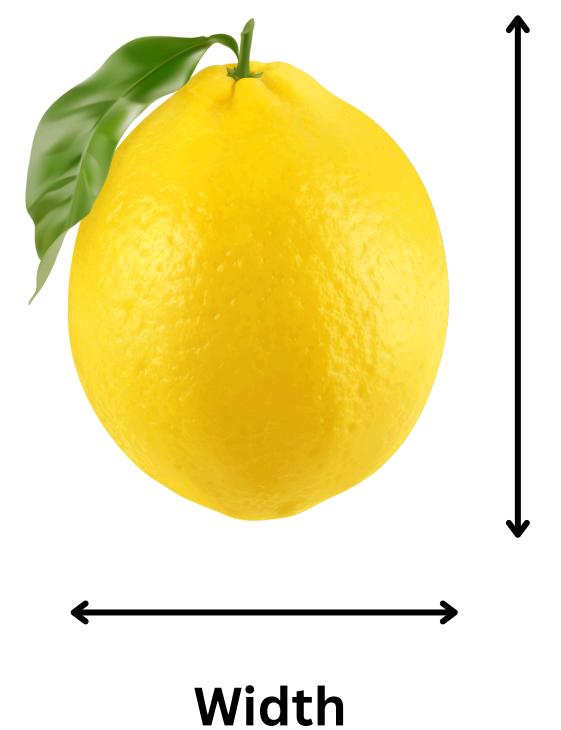
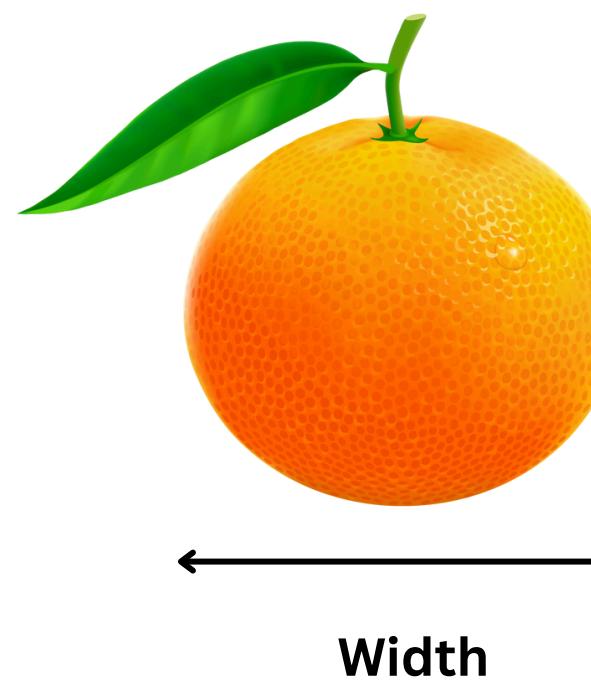
Define optimizing function (either a Loss or a Goal)

**Find Best  $\Theta$**

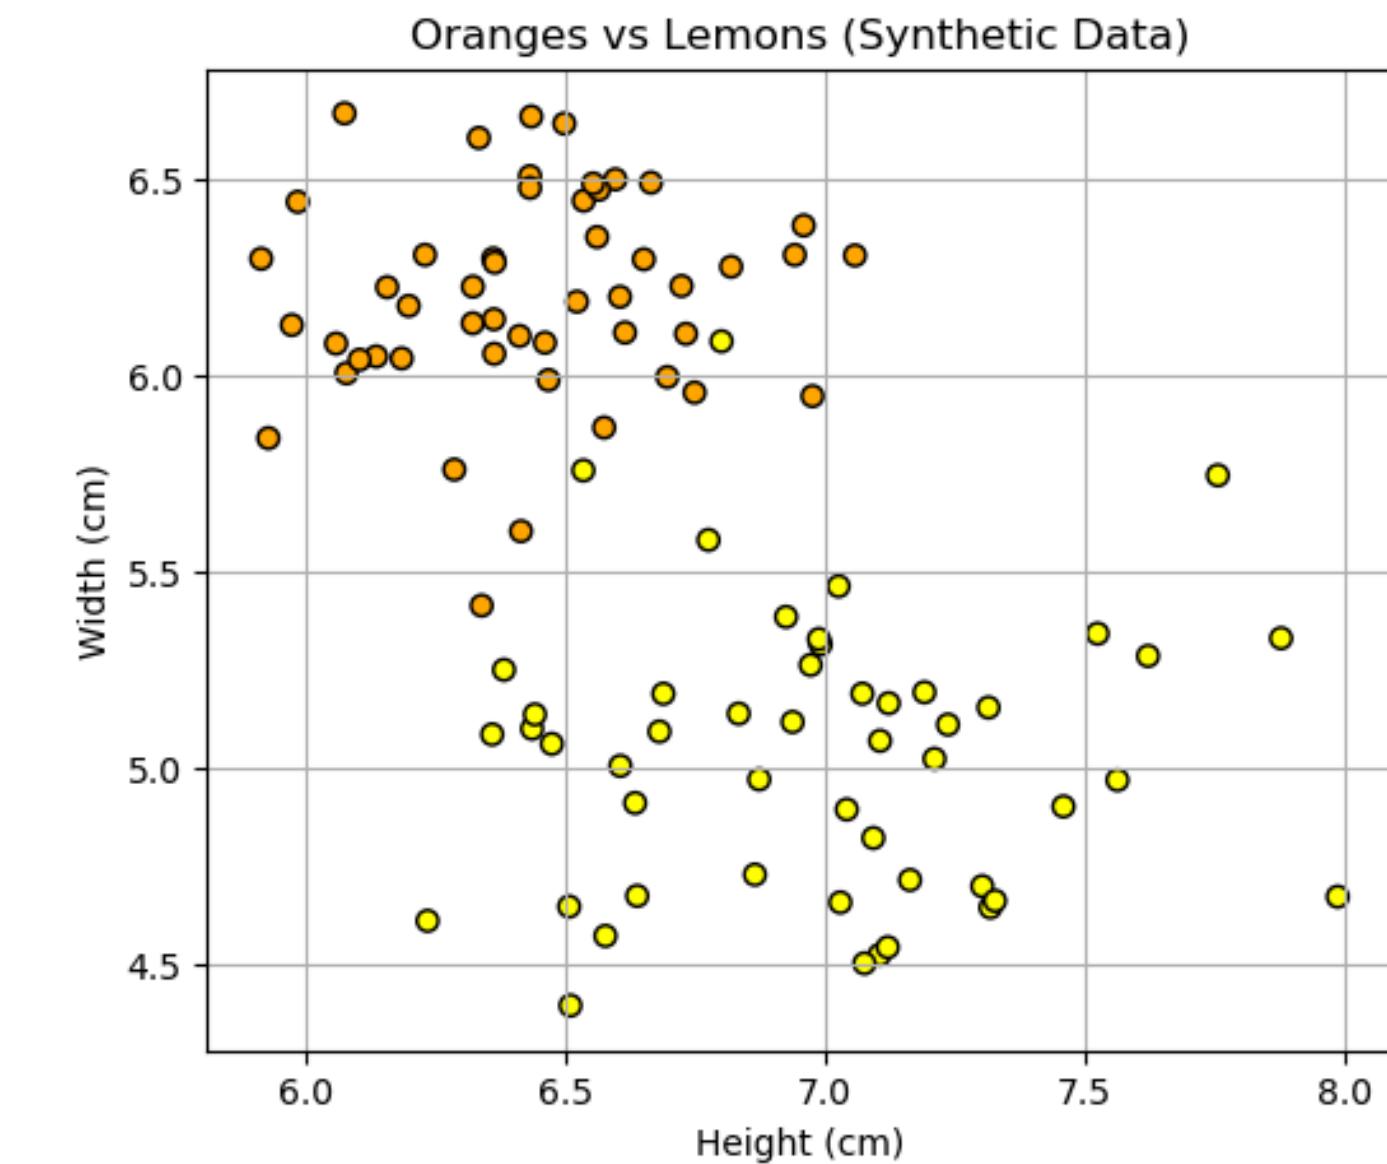
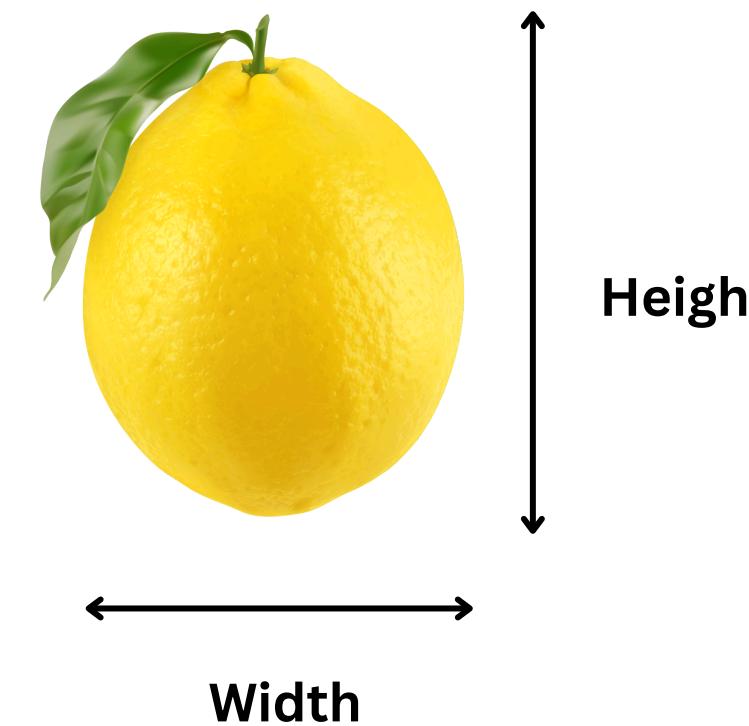
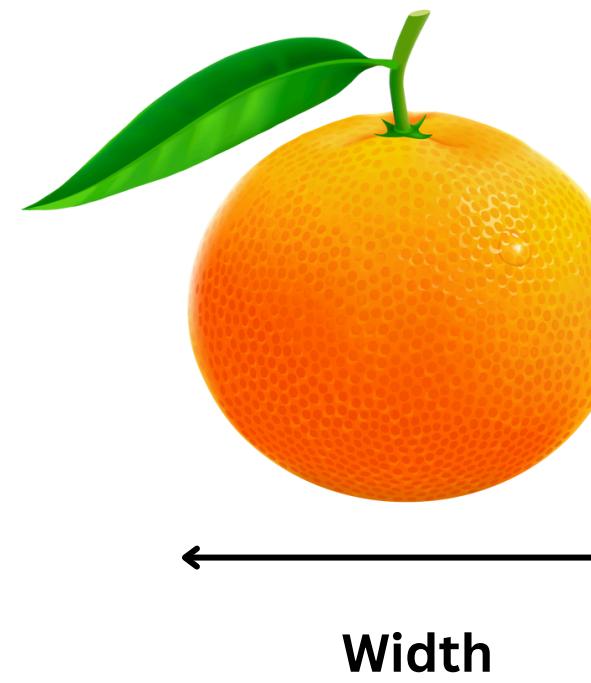


Optimization Procedure

# Practical use-case

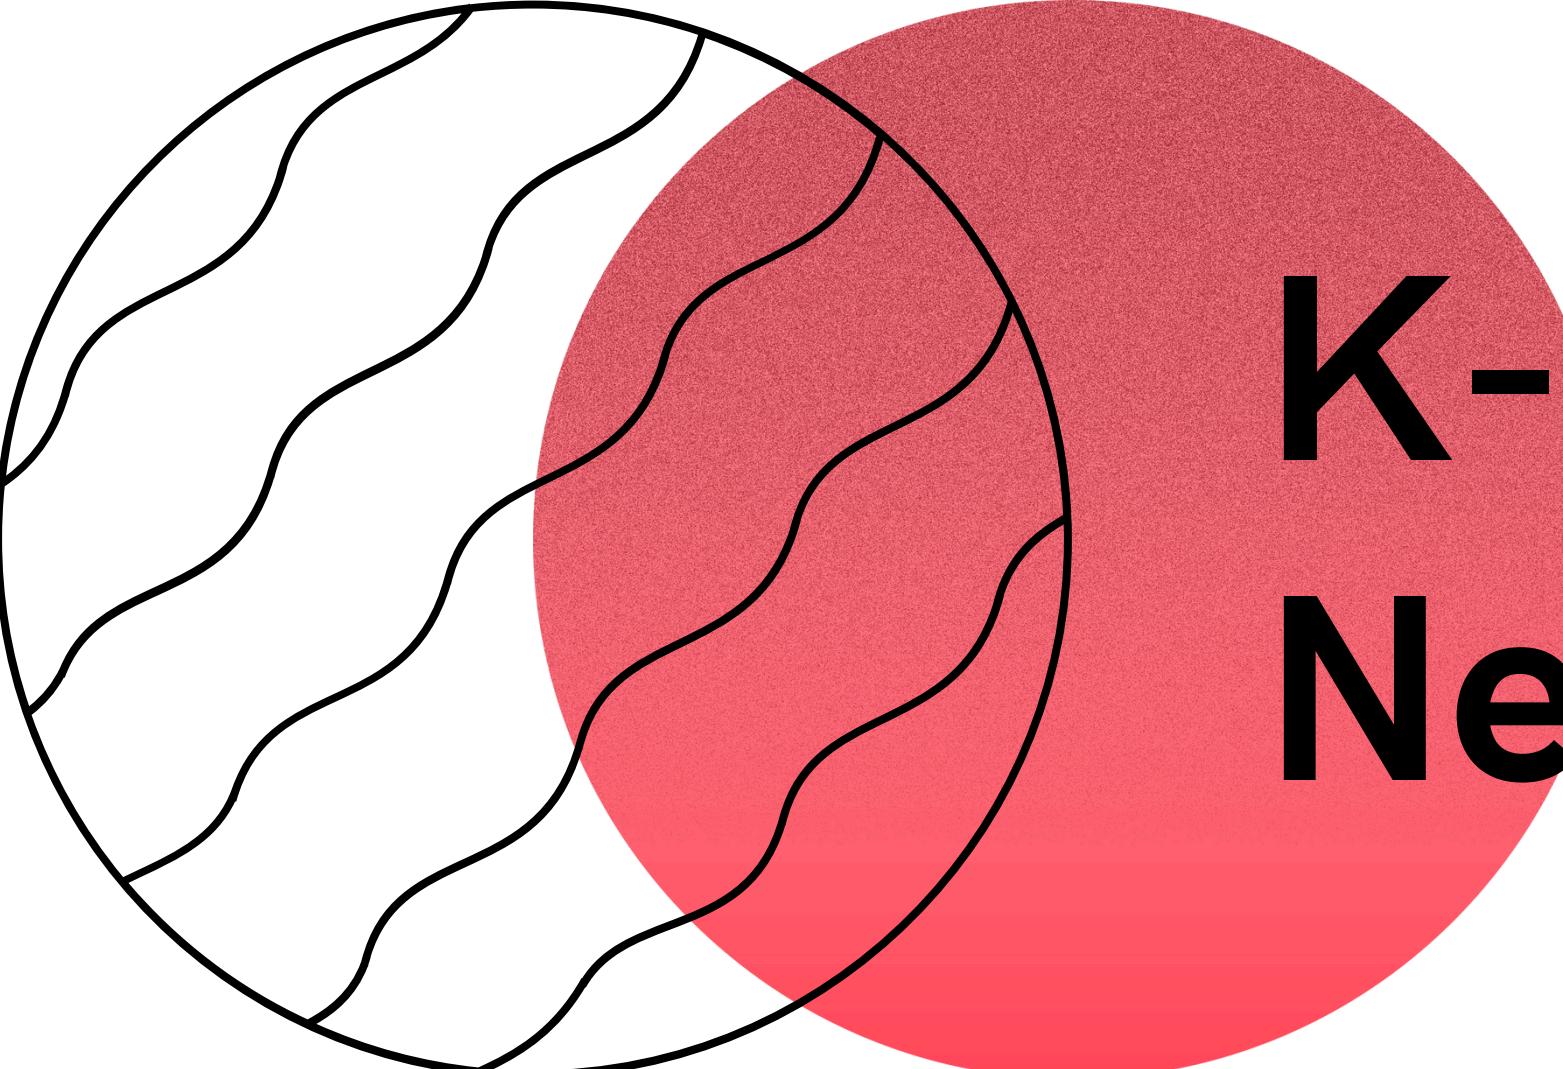


# Practical use-case



$f(\text{width}, \text{height}) : \mathbb{R}^2 \rightarrow \{0,1\}$ ,

0 - Orange  
1 - Lemon



# K-Nearest Neighbours

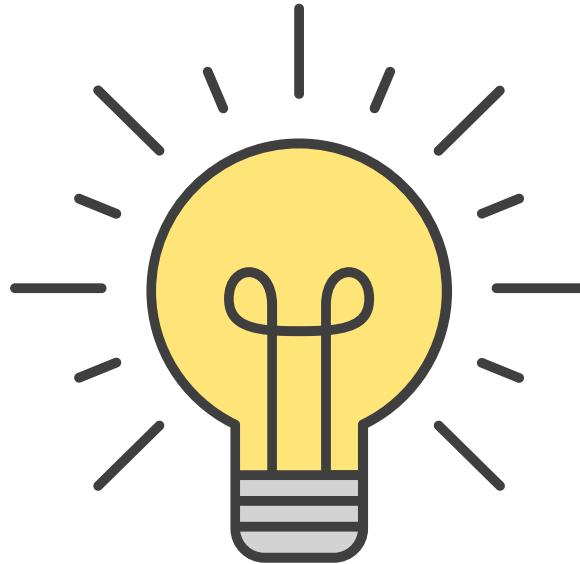
# K-NN



- **Non-parametric:** makes no assumptions about data distribution
- **Lazy learner:** defers computation to prediction time
- **Distance-based:** sensitive to feature scaling and irrelevant features

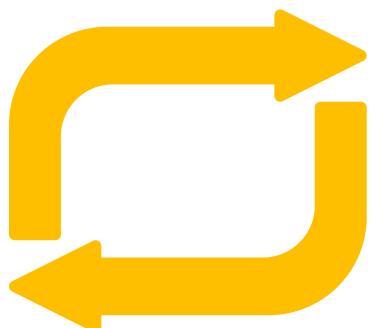
One of the simplest of all machine learning classifiers

# K-NN



## Idea:

- Similar examples have similar label.
- Classify new examples like similar training examples.



## Algorithm:

- Given some new example  $x$  for which we need to predict its class  $y$
- Find the  $k$  most similar training examples
- Classify  $x$  “like” these most similar examples (i.e. mode of target variable)

# K-NN

**k most similar instances:**

K Closest according to some distance measure

## Distance Functions

Euclidean

$$\sqrt{\sum_{i=1}^k (x_i - y_i)^2}$$

Manhattan

$$\sum_{i=1}^k |x_i - y_i|$$

Minkowski

$$\left( \sum_{i=1}^k (|x_i - y_i|)^q \right)^{1/q}$$

# K-NN

**k most similar instances:**

K Closest according to some distance measure

## Distance Functions

Euclidean

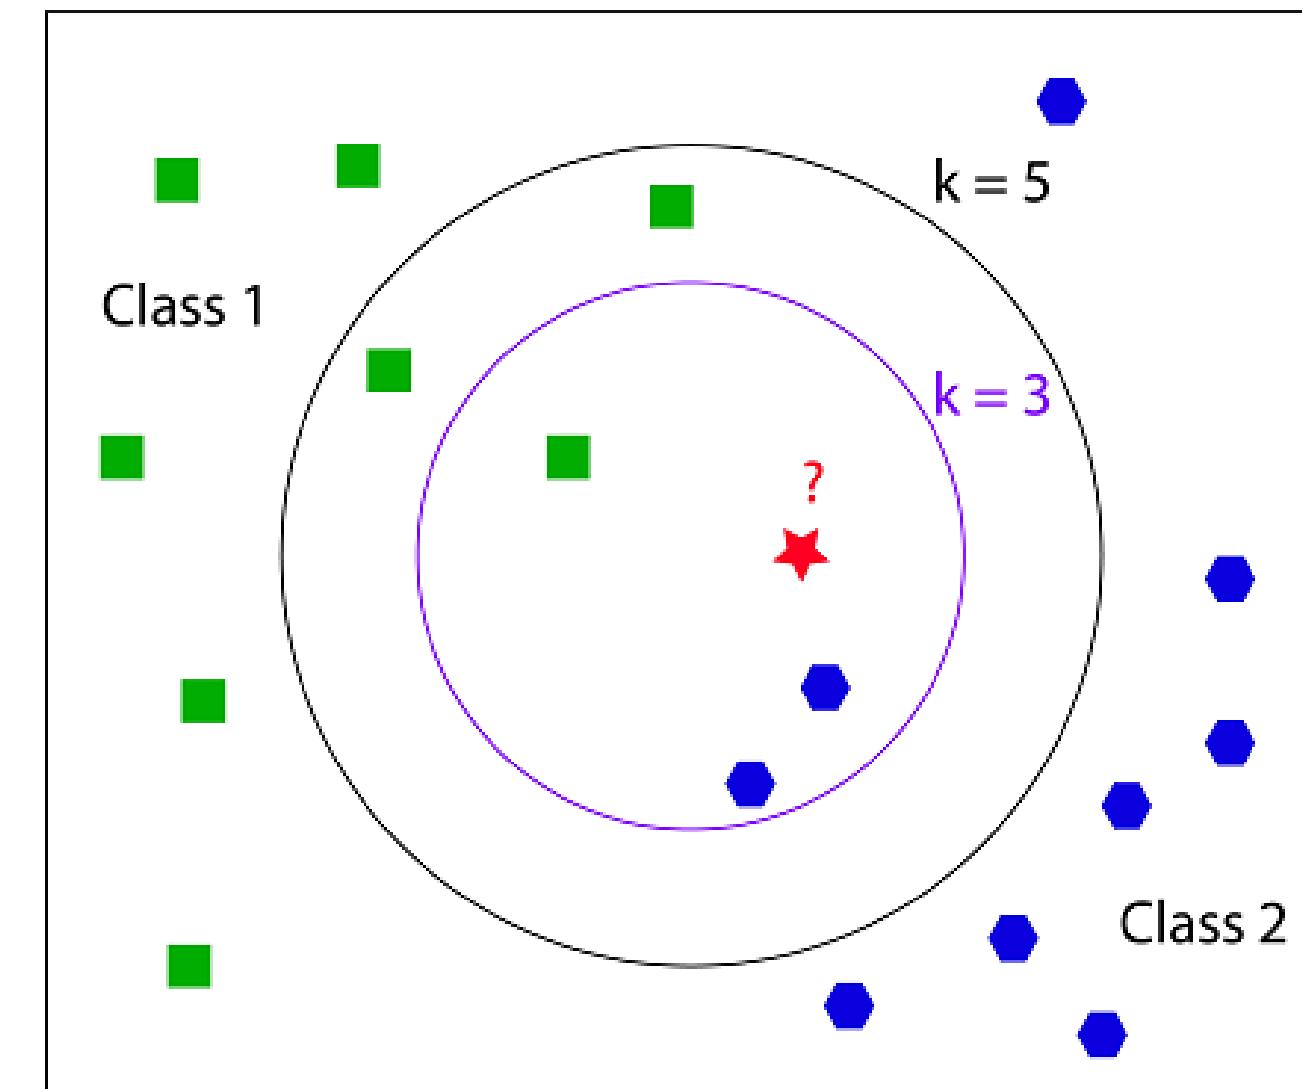
$$\sqrt{\sum_{i=1}^k (x_i - y_i)^2}$$

Manhattan

$$\sum_{i=1}^k |x_i - y_i|$$

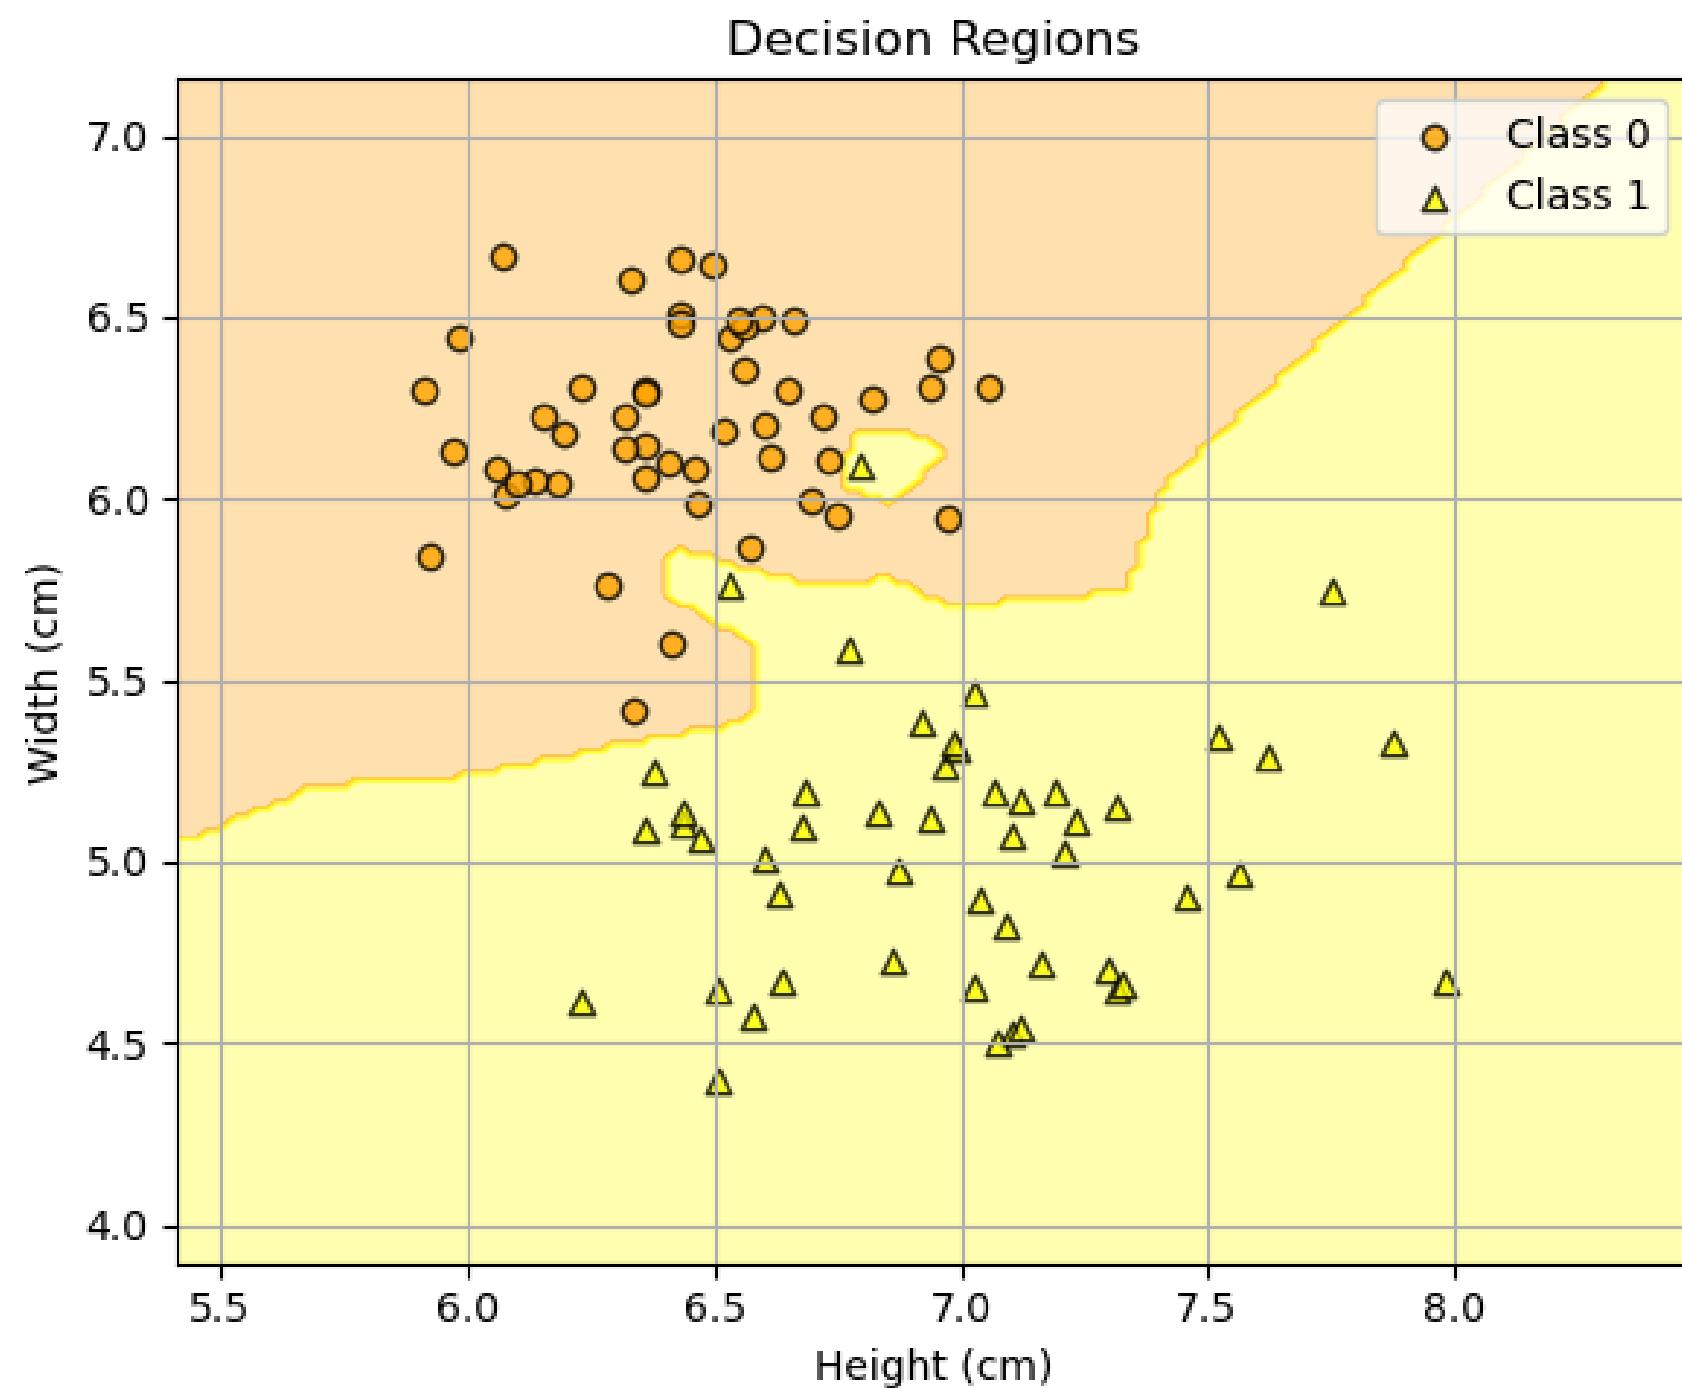
Minkowski

$$\left( \sum_{i=1}^k (|x_i - y_i|)^q \right)^{1/q}$$

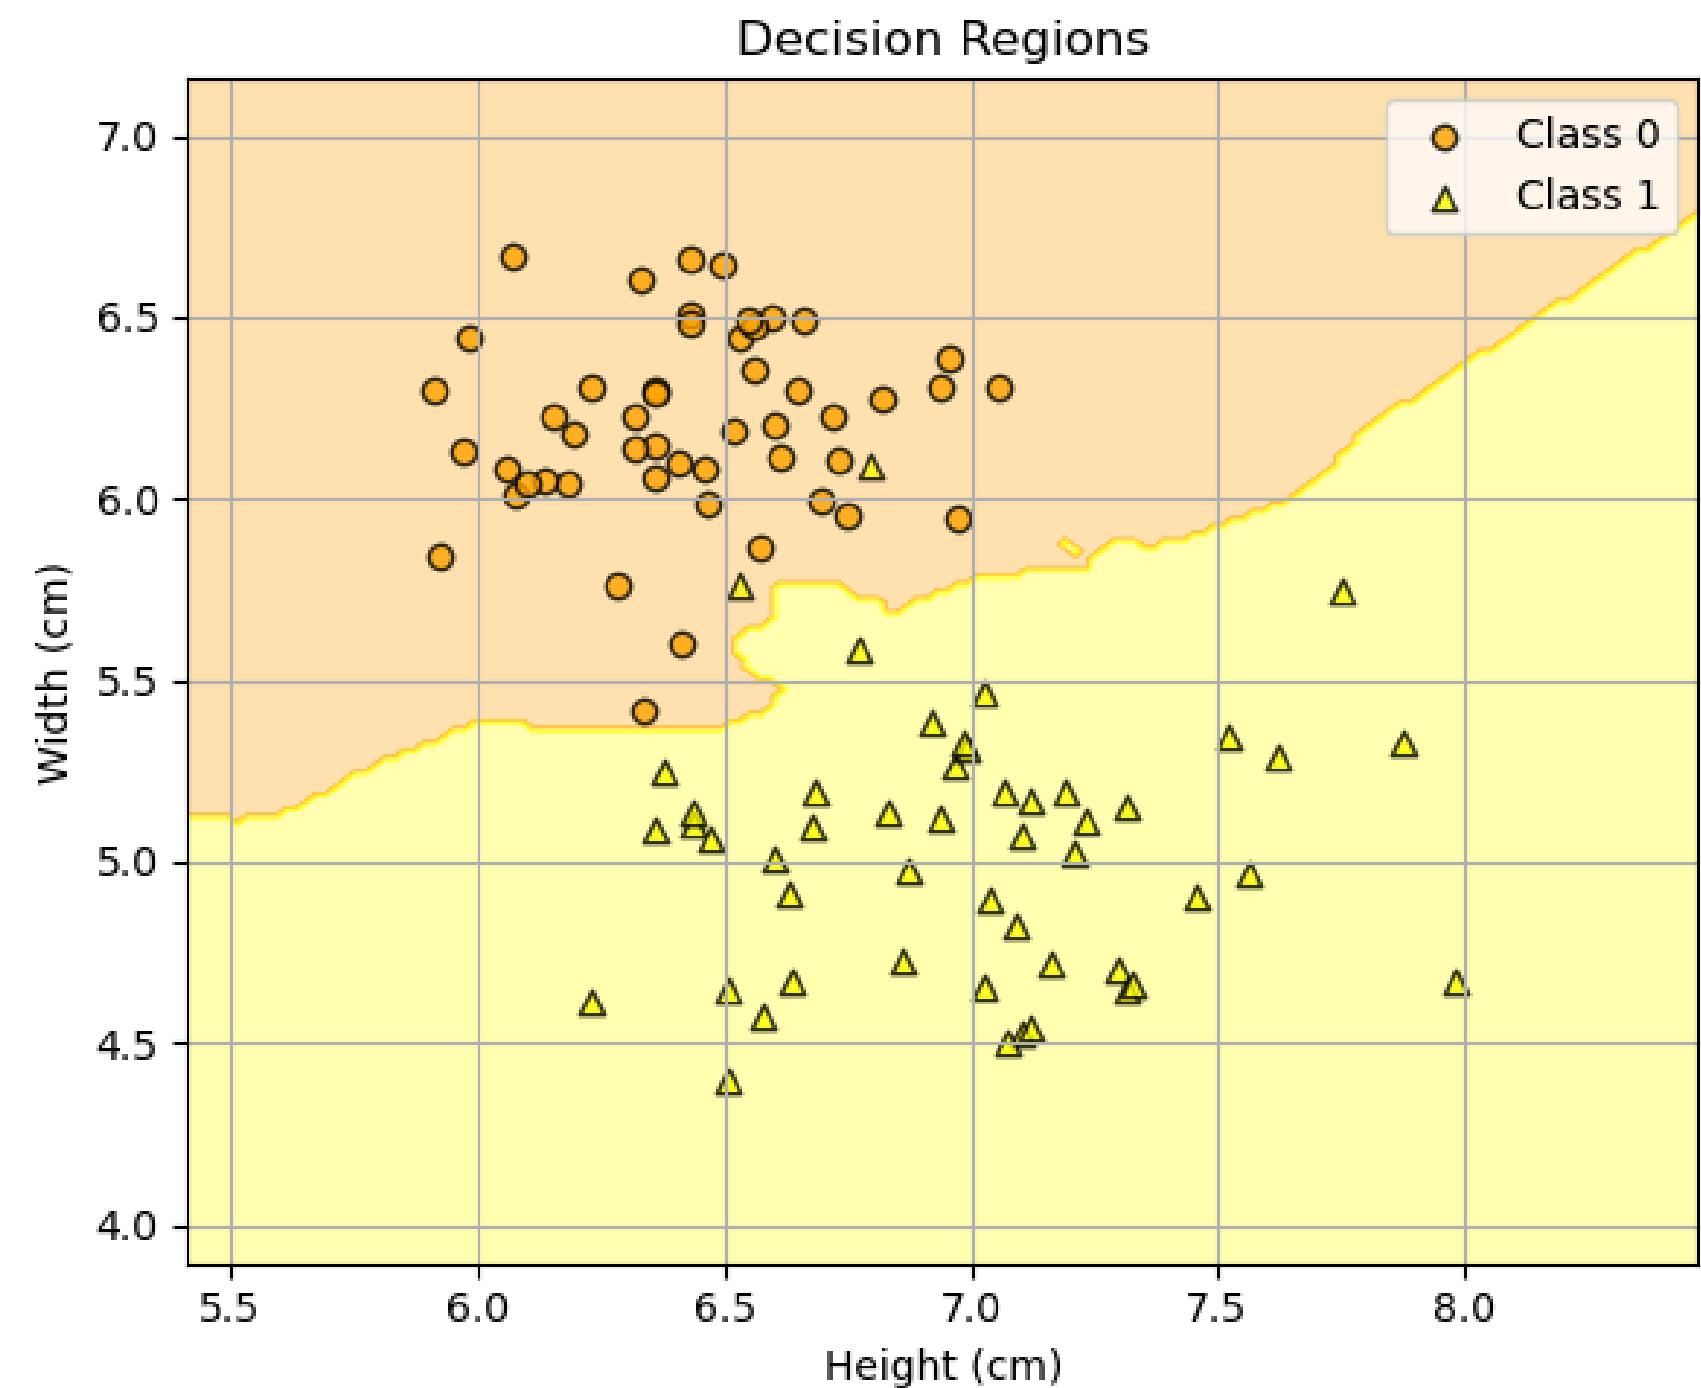


# K-NN

1-NN



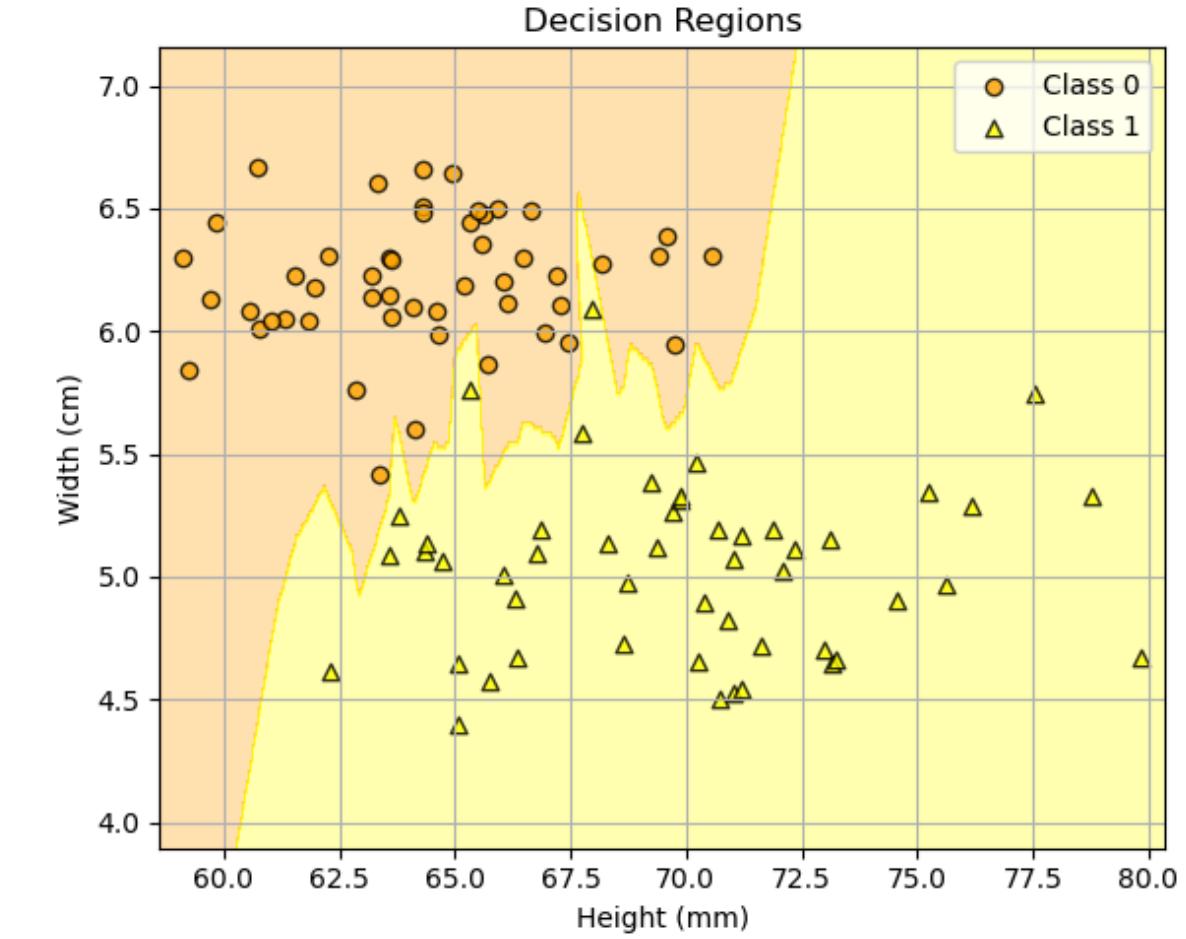
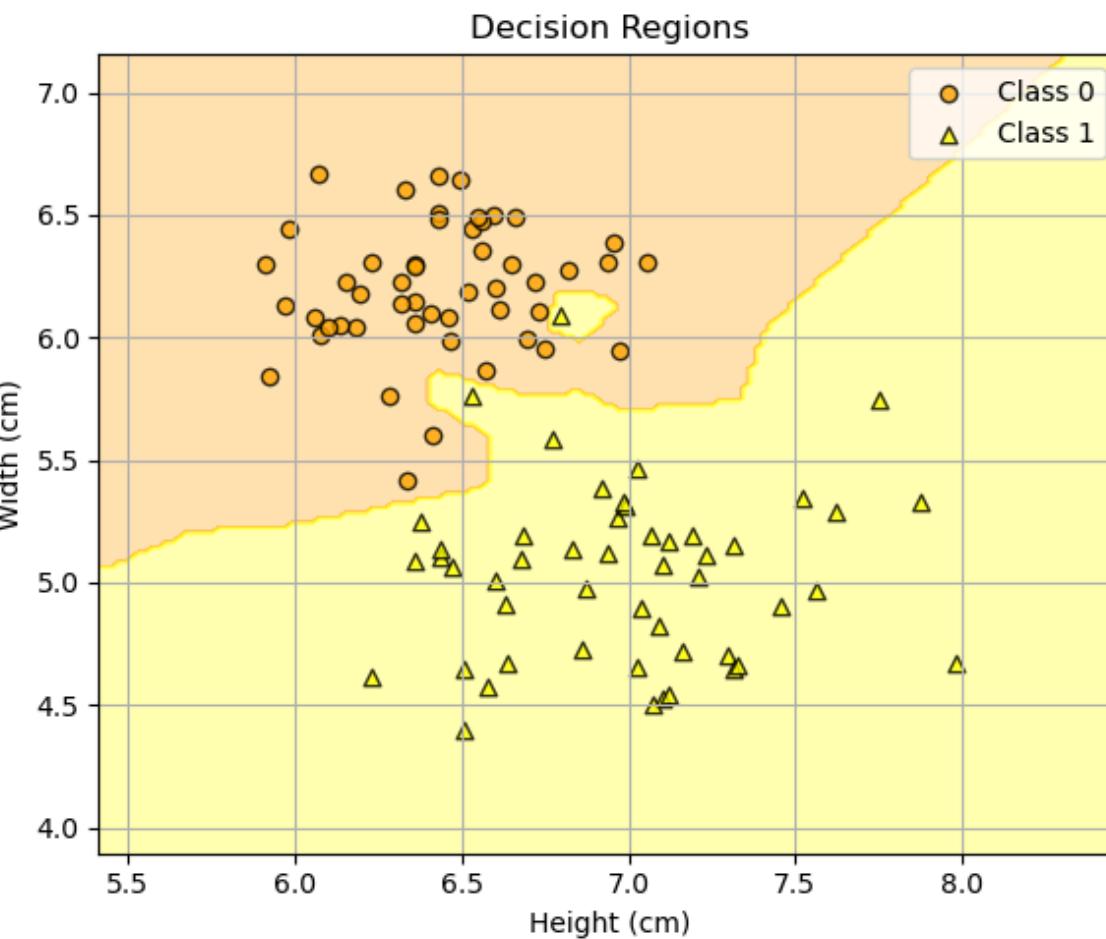
3-NN



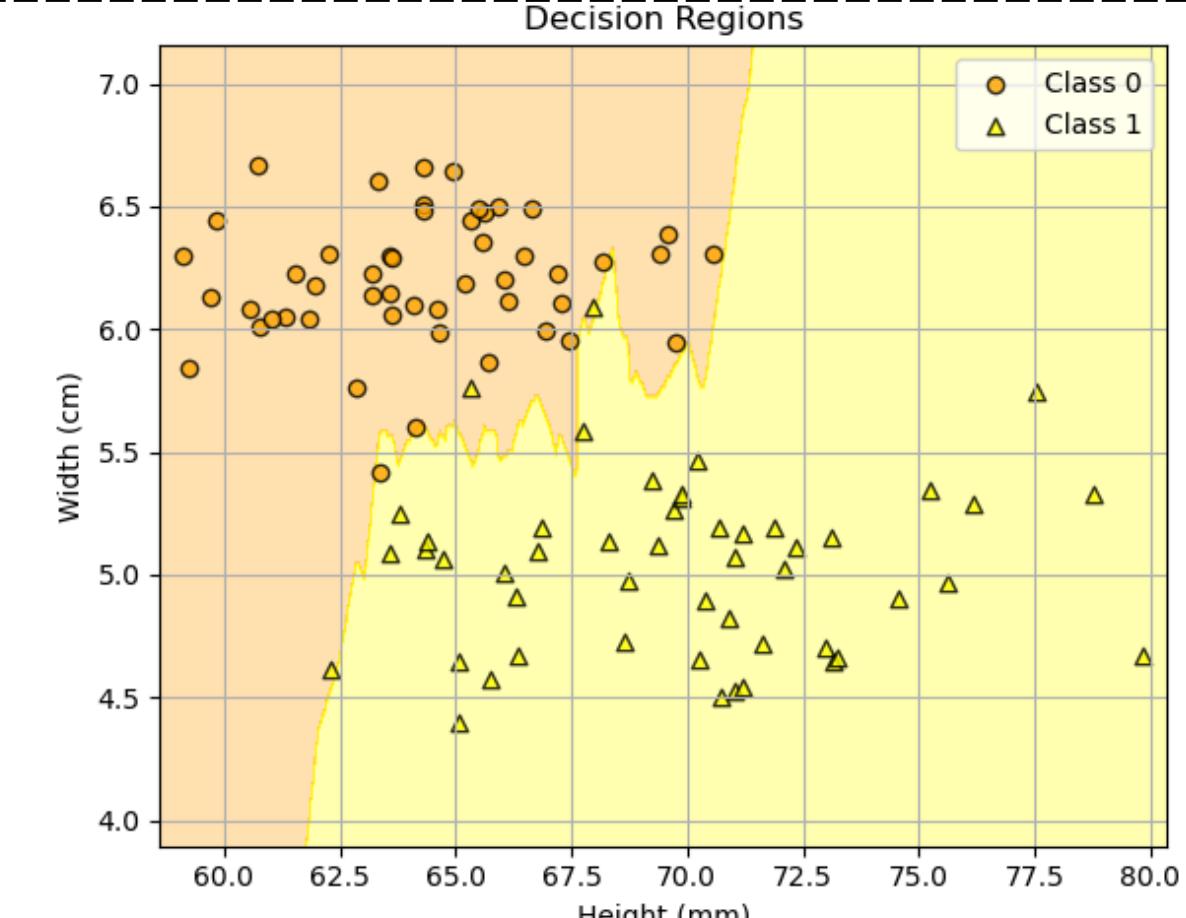
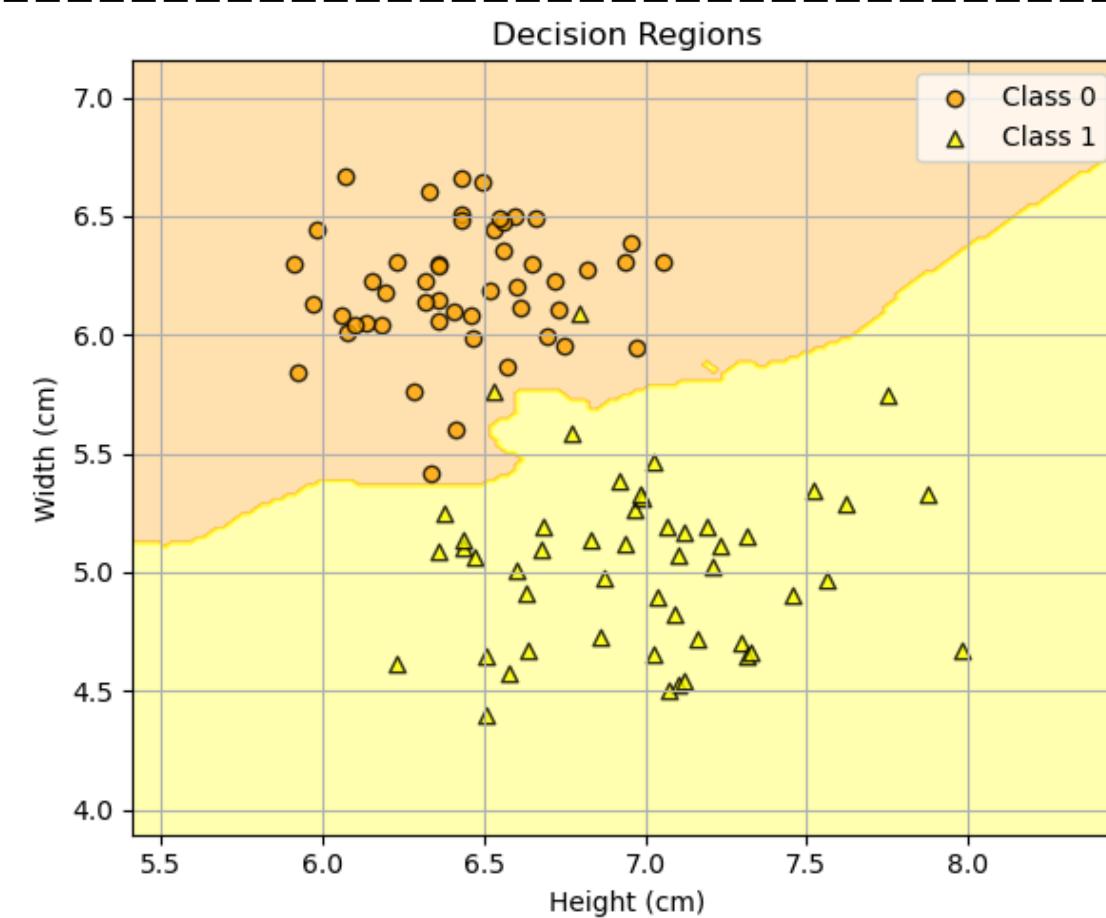
# K-NN

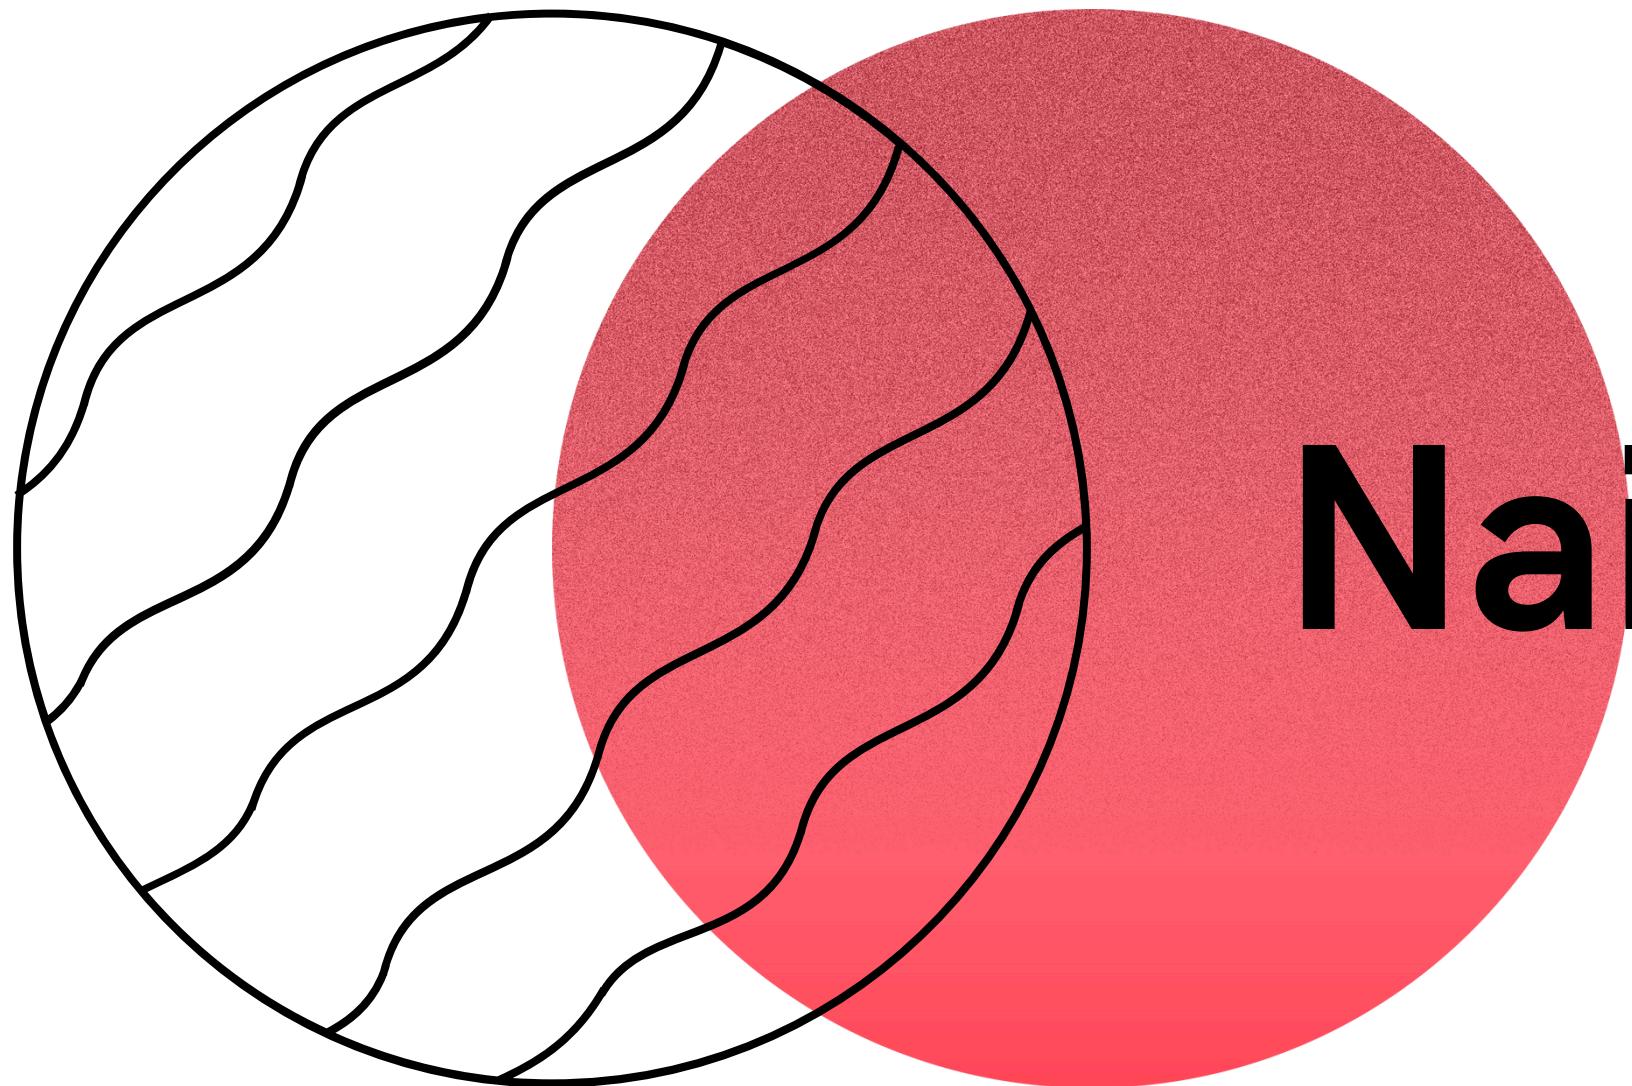
## Importance of Scale

1-NN



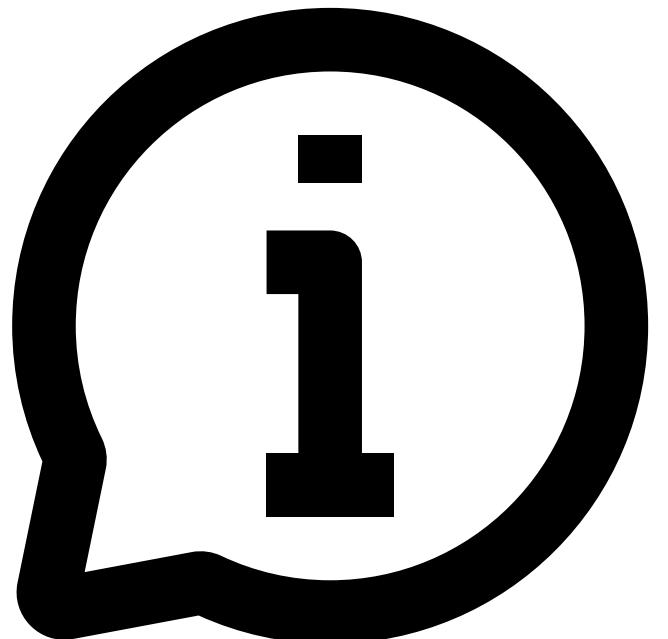
3-NN





# Naive Bayes

# Naive Bayes



- **Probabilistic**

For a pair of (height,width) what is more likely?  
Lemon or Orange?

$$P(y = 1 | (\text{height}, \text{width})) \quad (> \text{ or } < ?) \quad P(y = 0 | (\text{height}, \text{width}))$$

- **Naive:** Assumes predictor independance

# Naive Bayes



Idea:

$$P(c|x) = \frac{P(x|c)P(c)}{P(x)}$$

Diagram illustrating the components of the Naive Bayes formula:

- Likelihood:  $P(x|c)$  (indicated by a blue arrow pointing to the numerator term  $P(x|c)$ )
- Class Prior Probability:  $P(c)$  (indicated by a blue arrow pointing to the numerator term  $P(c)$ )
- Posterior Probability:  $P(c|x)$  (indicated by a blue arrow pointing to the entire fraction  $P(c|x)$ )
- Predictor Prior Probability:  $P(x)$  (indicated by a blue arrow pointing to the denominator term  $P(x)$ )

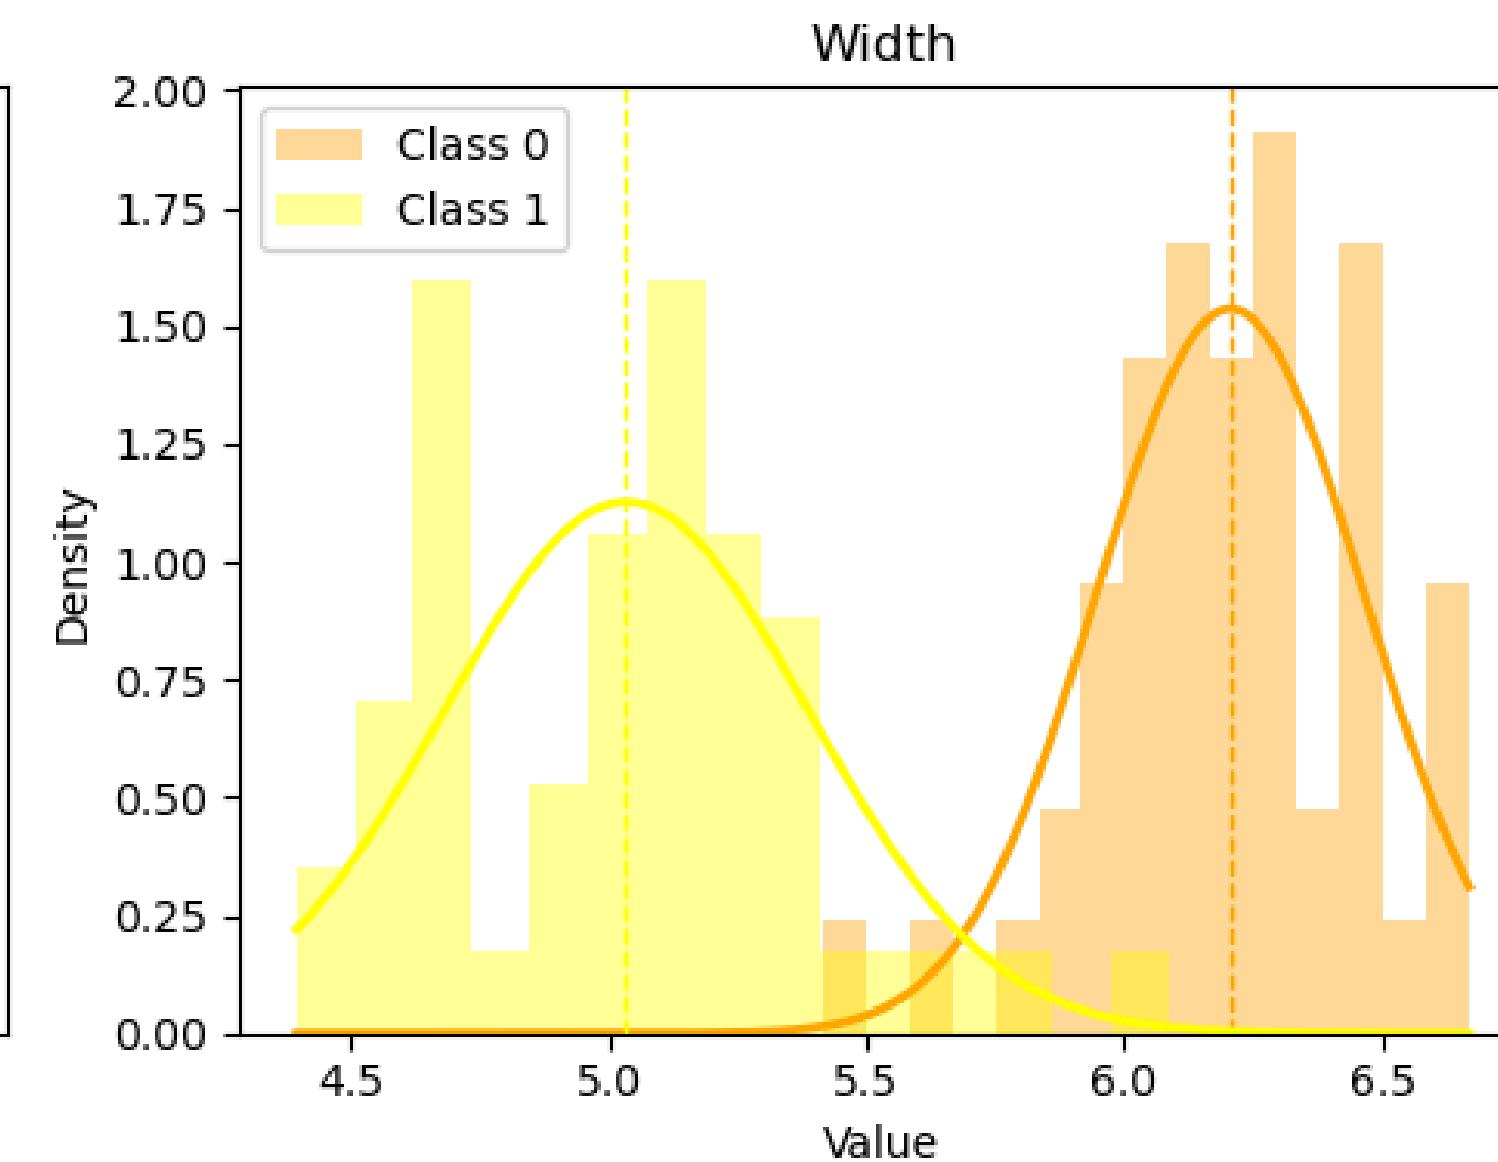
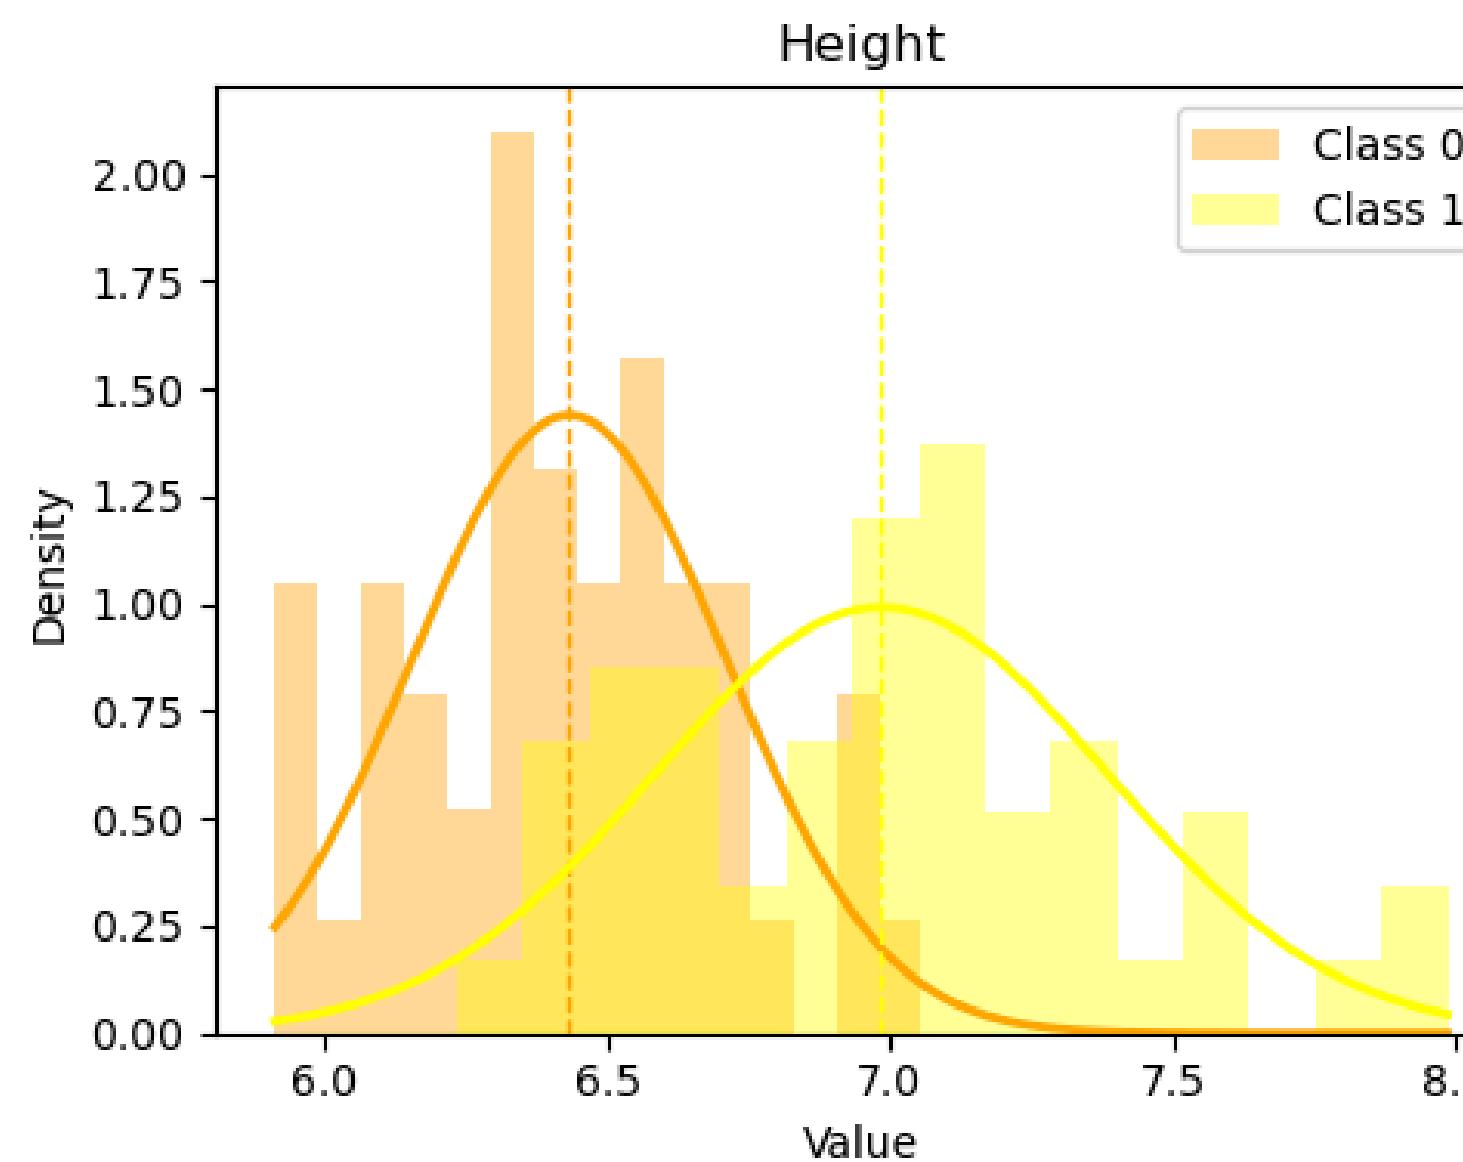
$$P(c|x) \propto P(x_1|c) \times P(x_2|c) \times \cdots \times P(x_n|c) \times P(c)$$

Assume a distribution for  $P(x|c)$ , and use parametric statistics to find distribution parameters

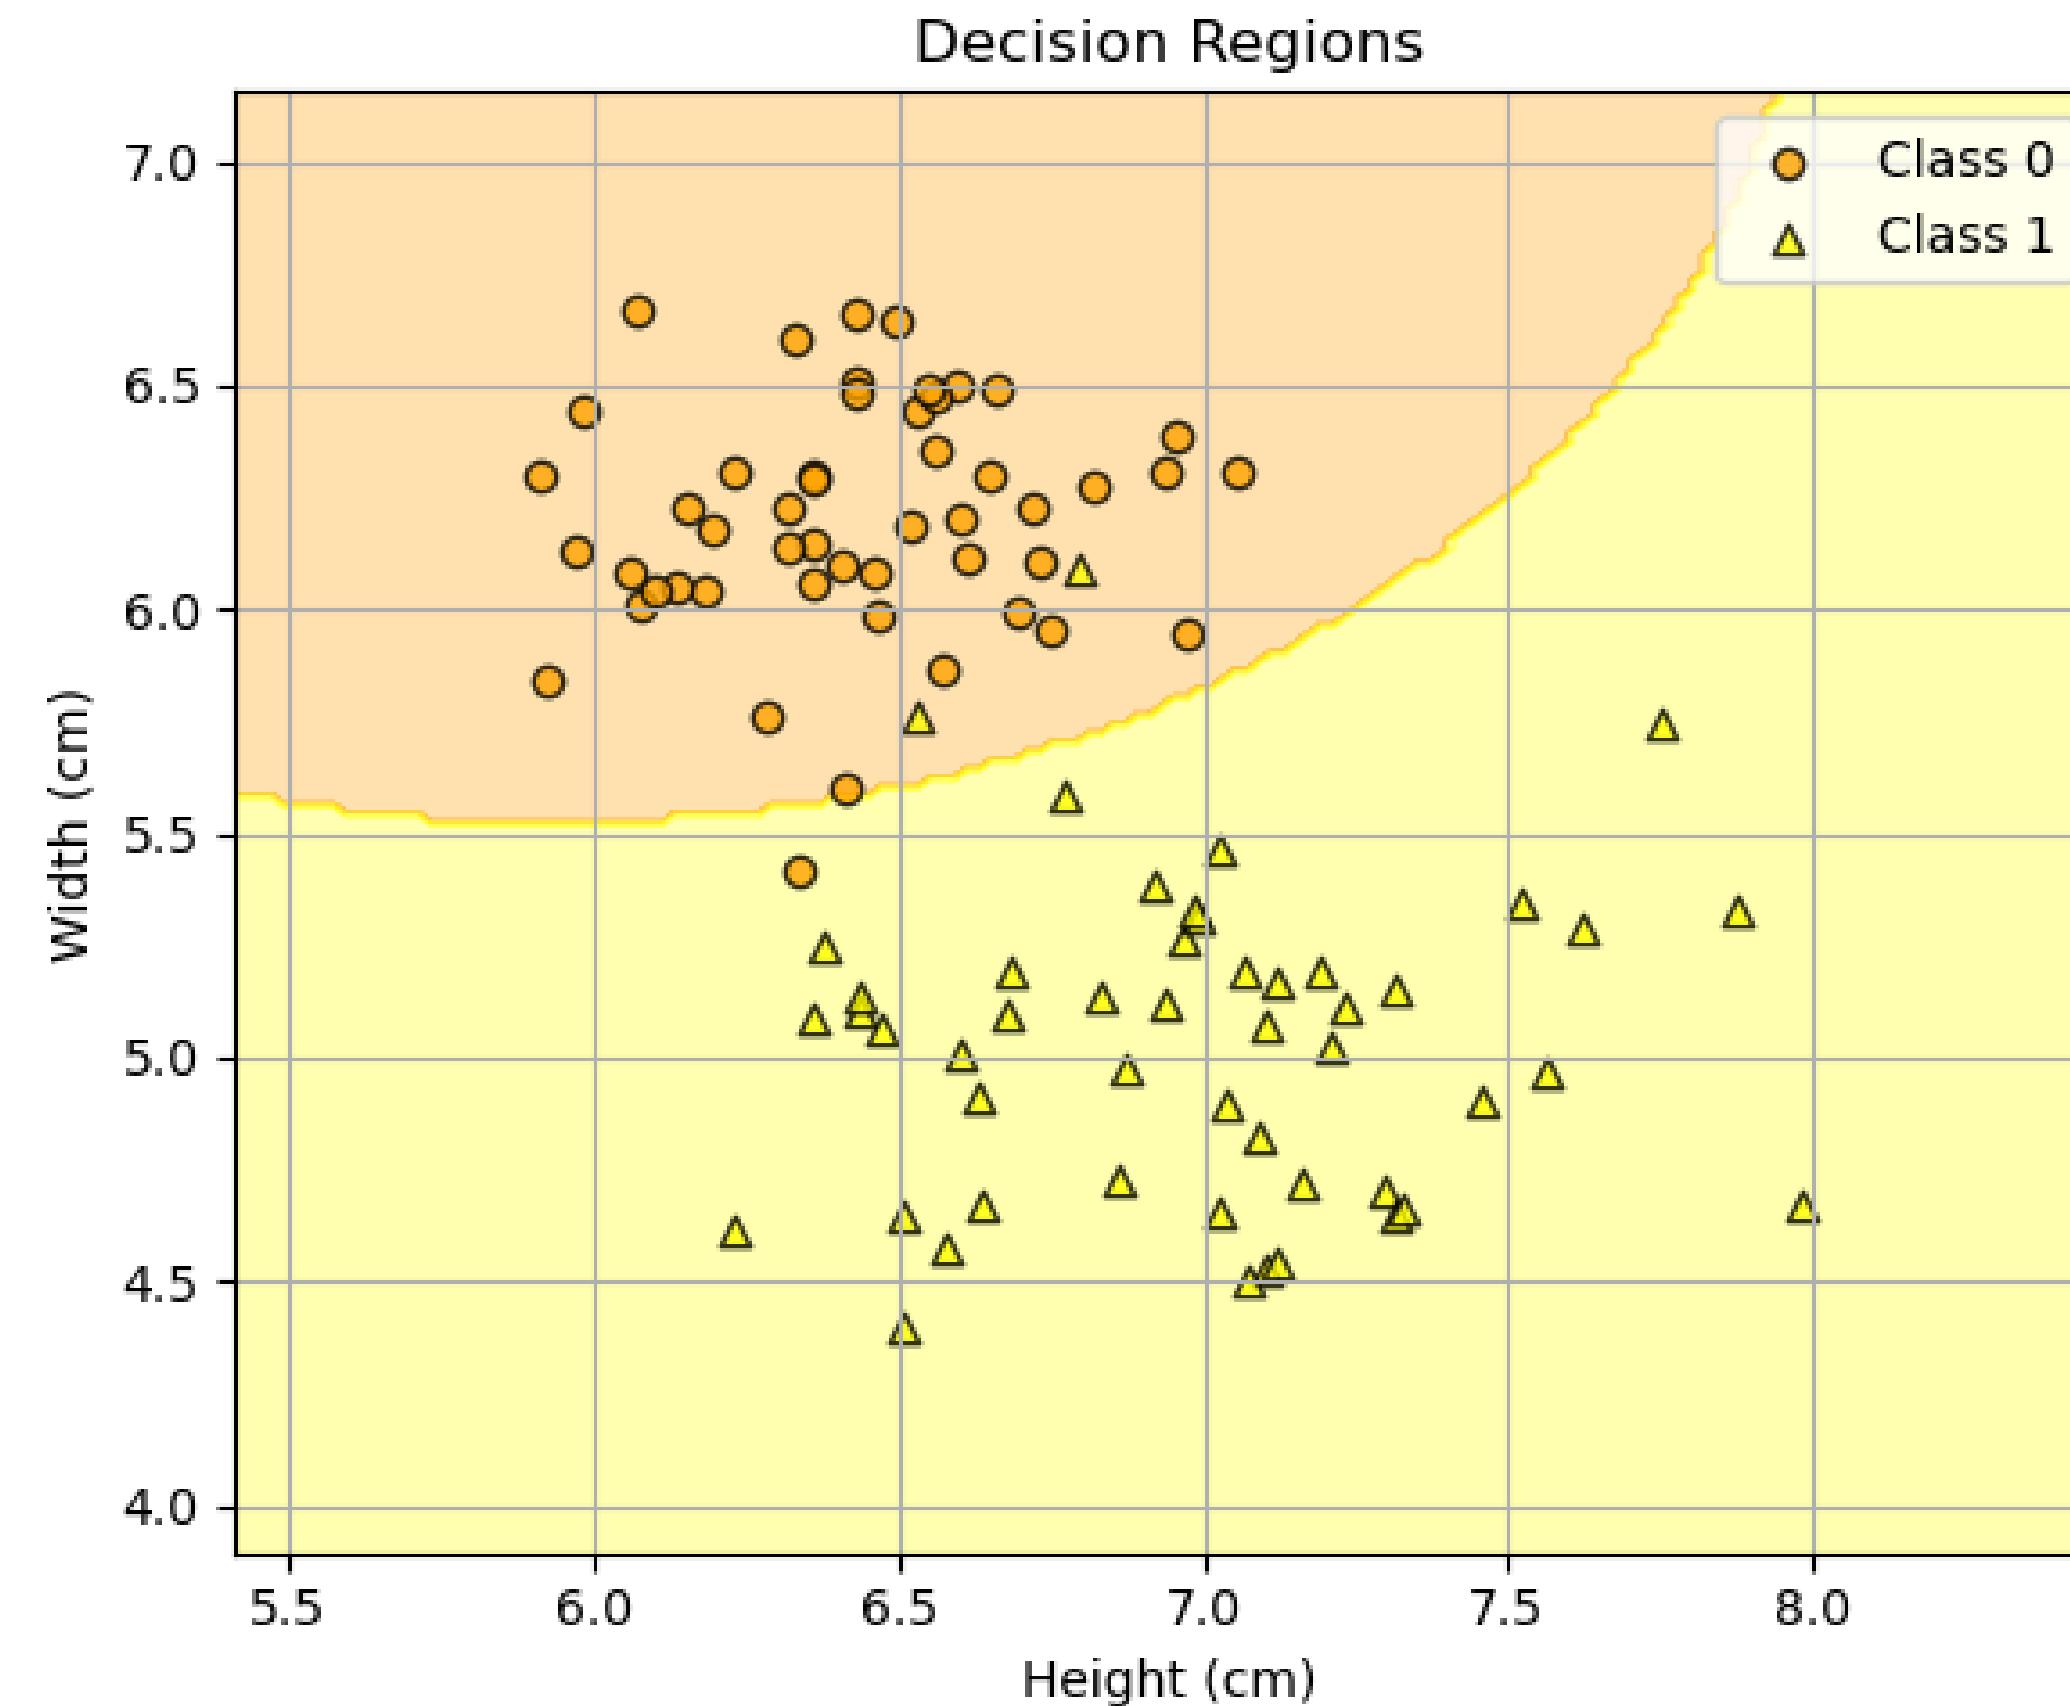
# Naive Bayes

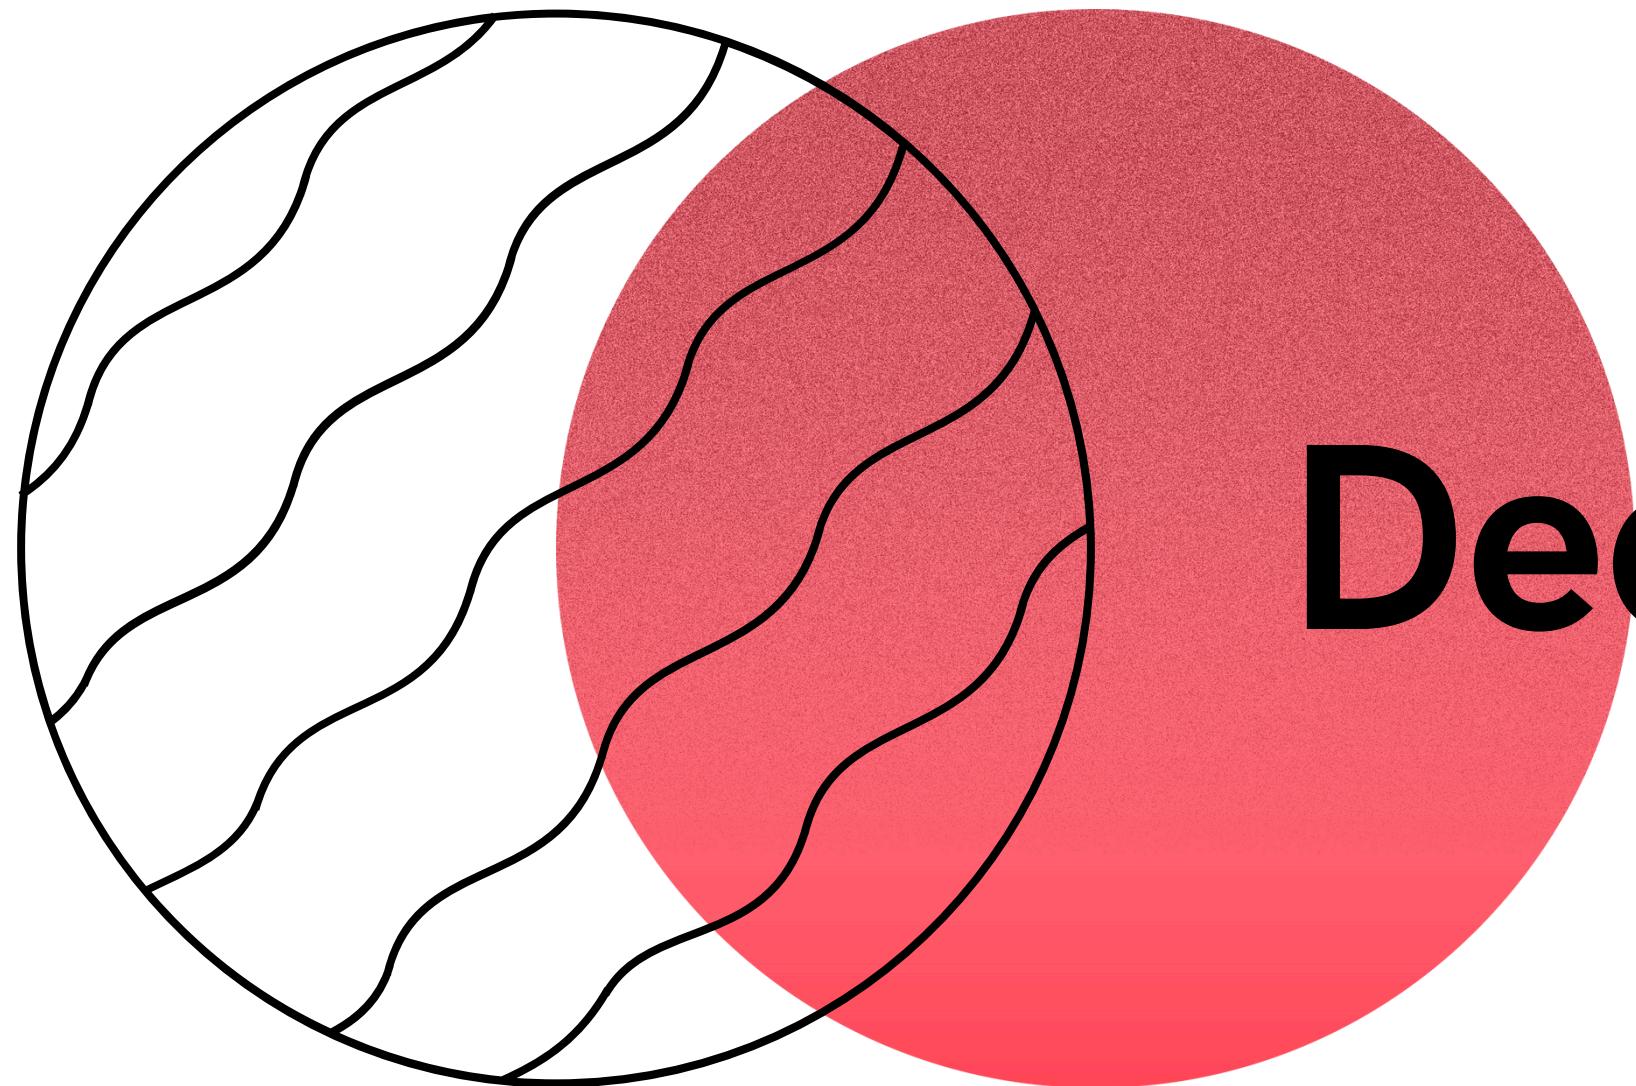
$$\mathcal{N}(\mu, \sigma^2)$$

$$\frac{1}{\sqrt{2\pi}\sigma^2} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$



# Naive Bayes





# Decision Trees

# Decision Trees

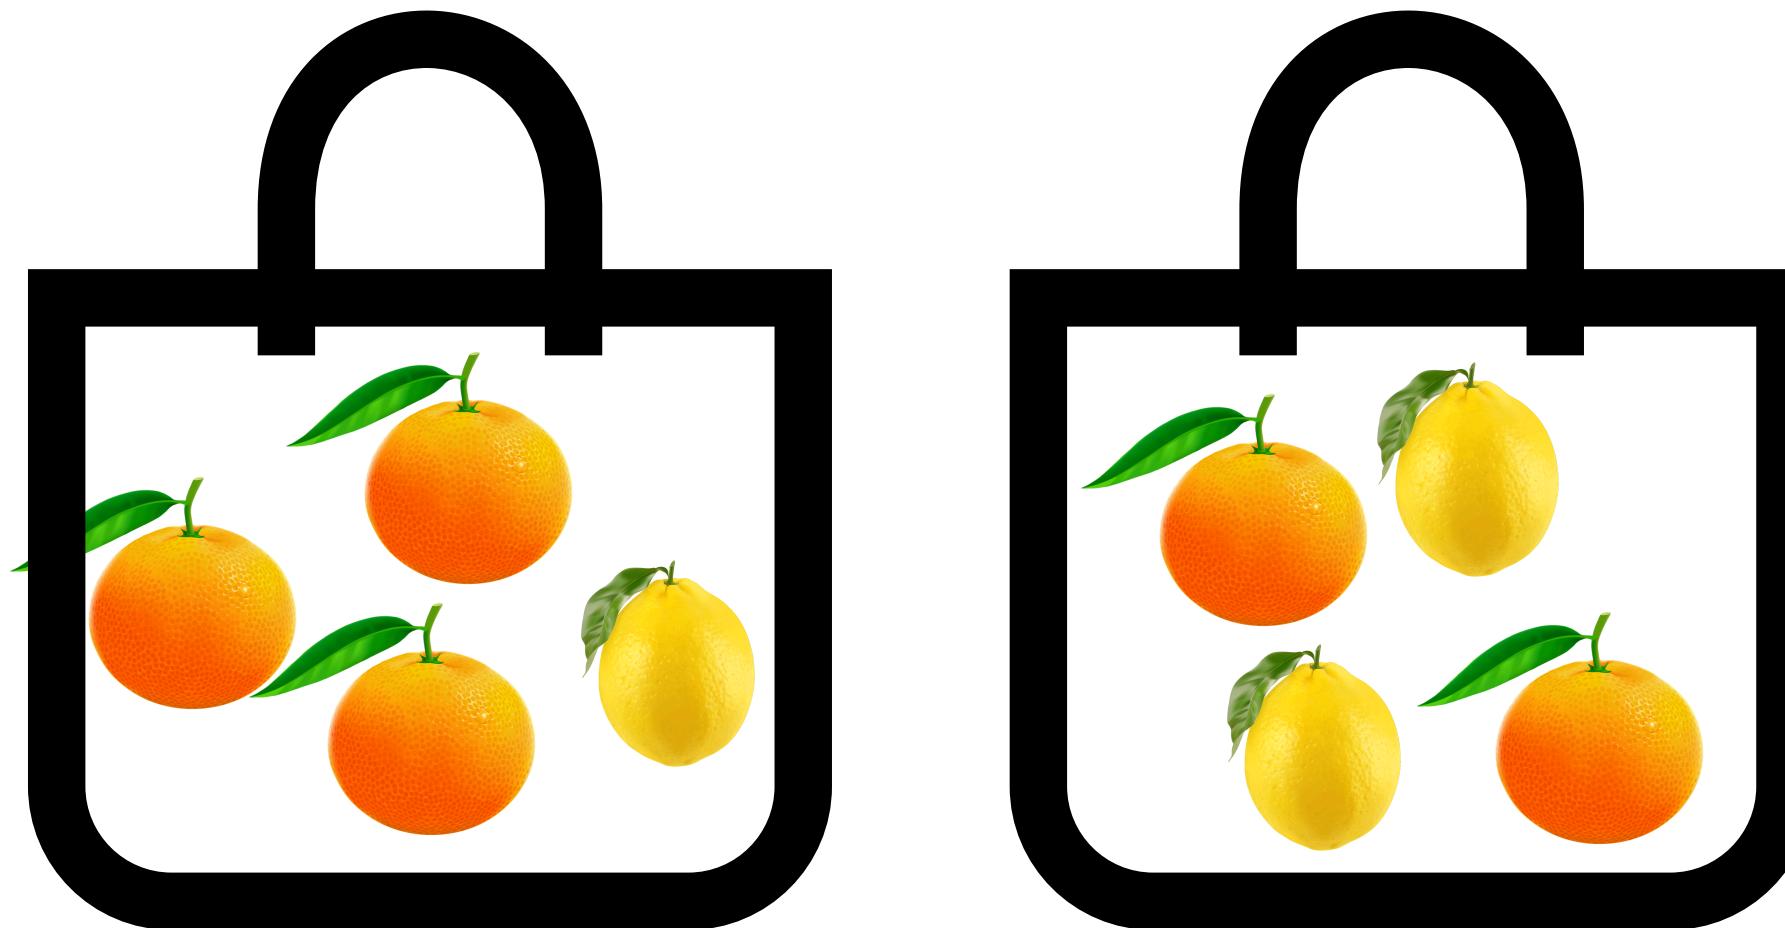


- Mimics human decision-making by learning **simple if-then rules** from the data
- Builds a tree-like model of decisions:
  - **Internal nodes:** feature-based decisions
  - **Branches:** outcomes of the decision
  - **Leaves:** final predictions
- **Greedy optimization strategy**

# Decision Trees



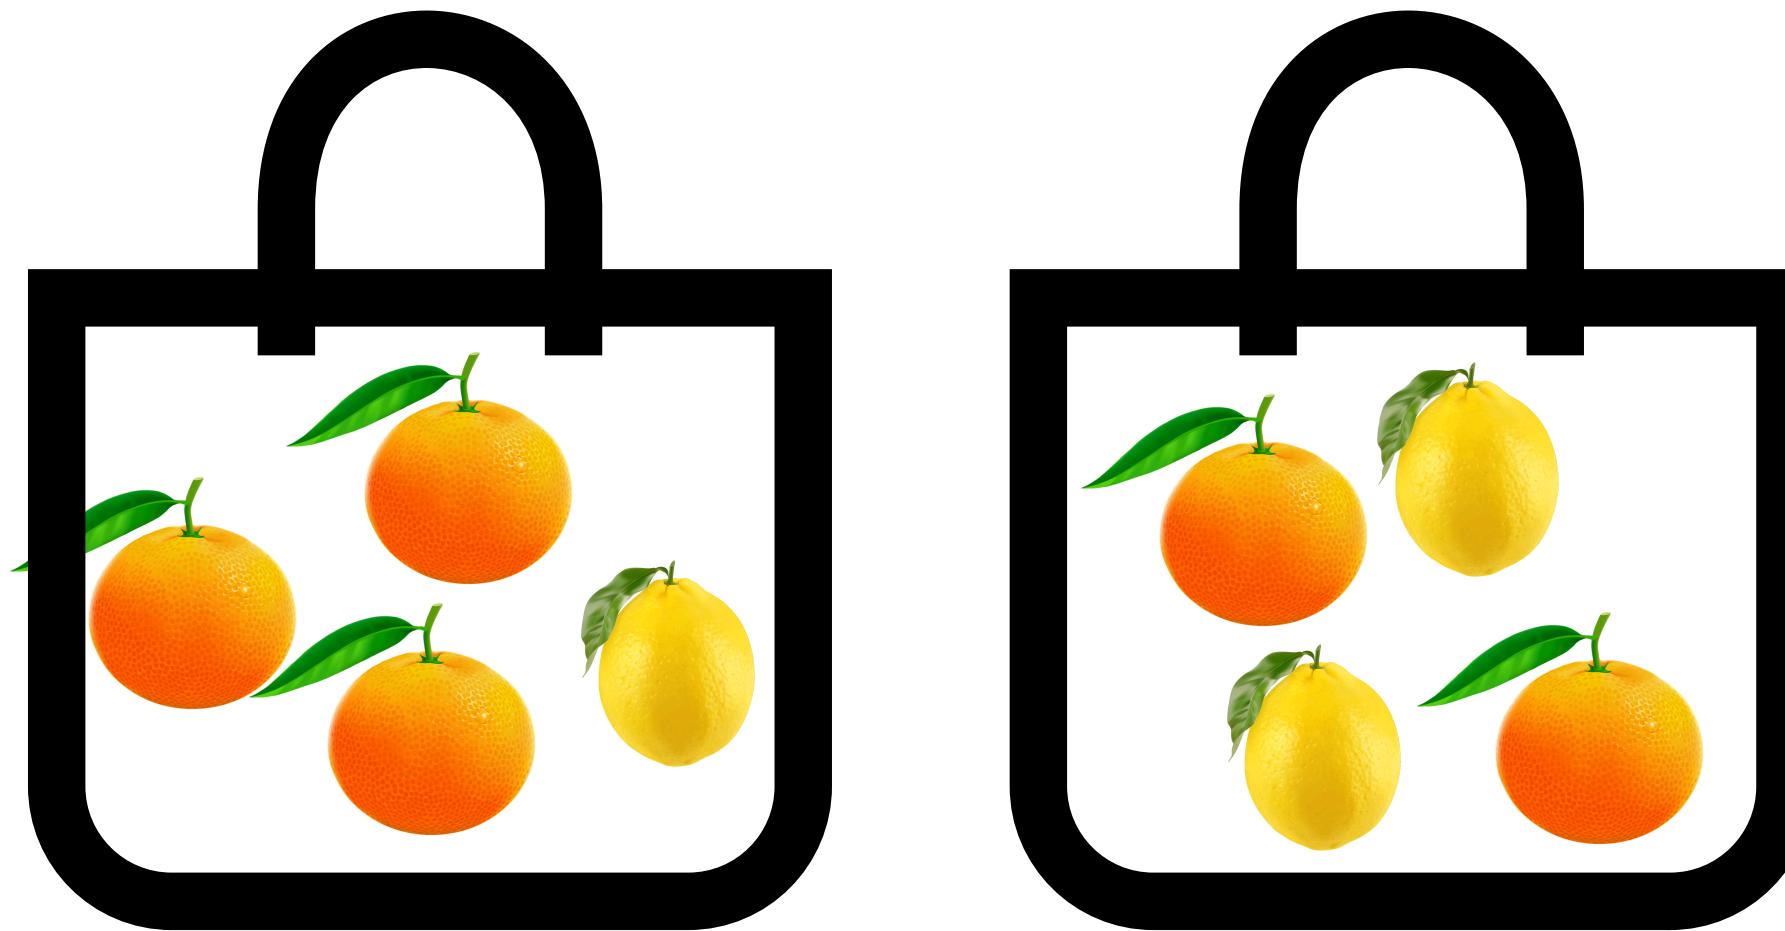
**Idea:** In which bag can u make a more confident decision of what fruit you'll randomly take?



# Decision Trees



**Idea:** In which bag can u make a more confident decision of what fruit you'll randomly take?



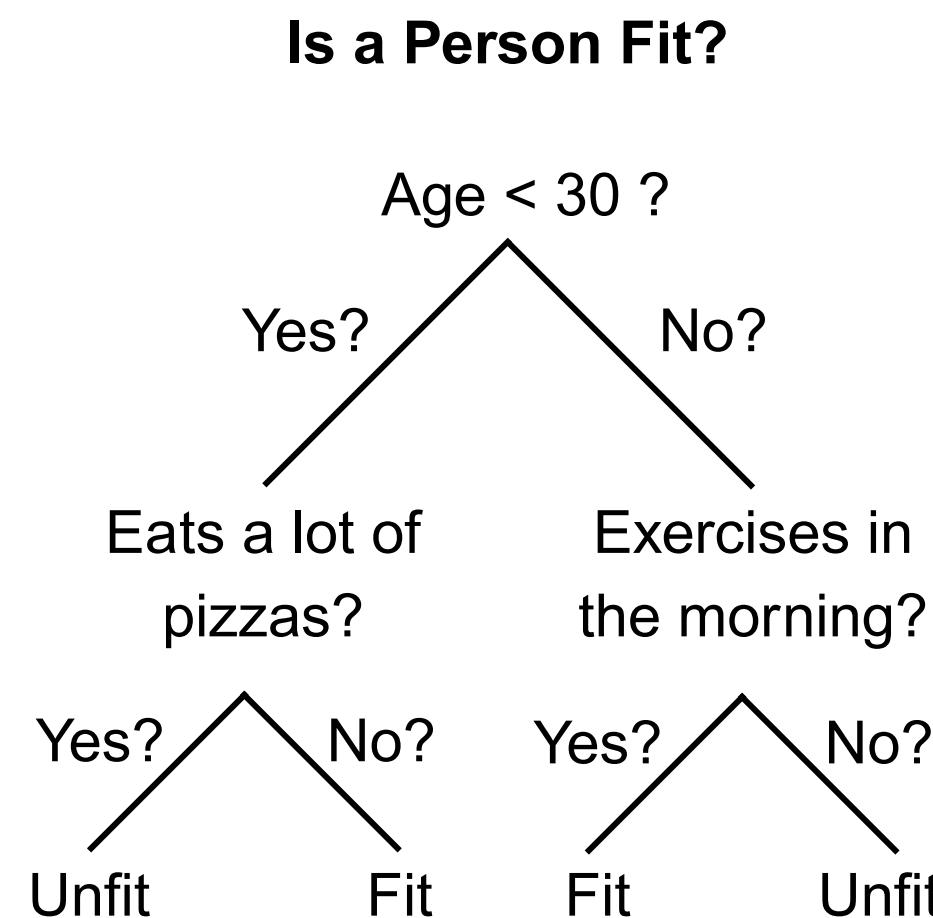
The right bag is more ‘mixed’

# Decision Trees

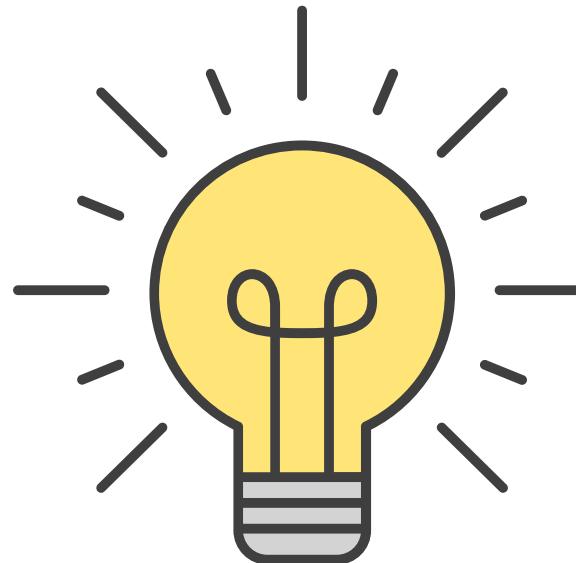


## Idea:

- Create disjoint regions in the space, where data is less ‘mixed’ so decisions locally can be more confident.
- Done by creating if-then questions!



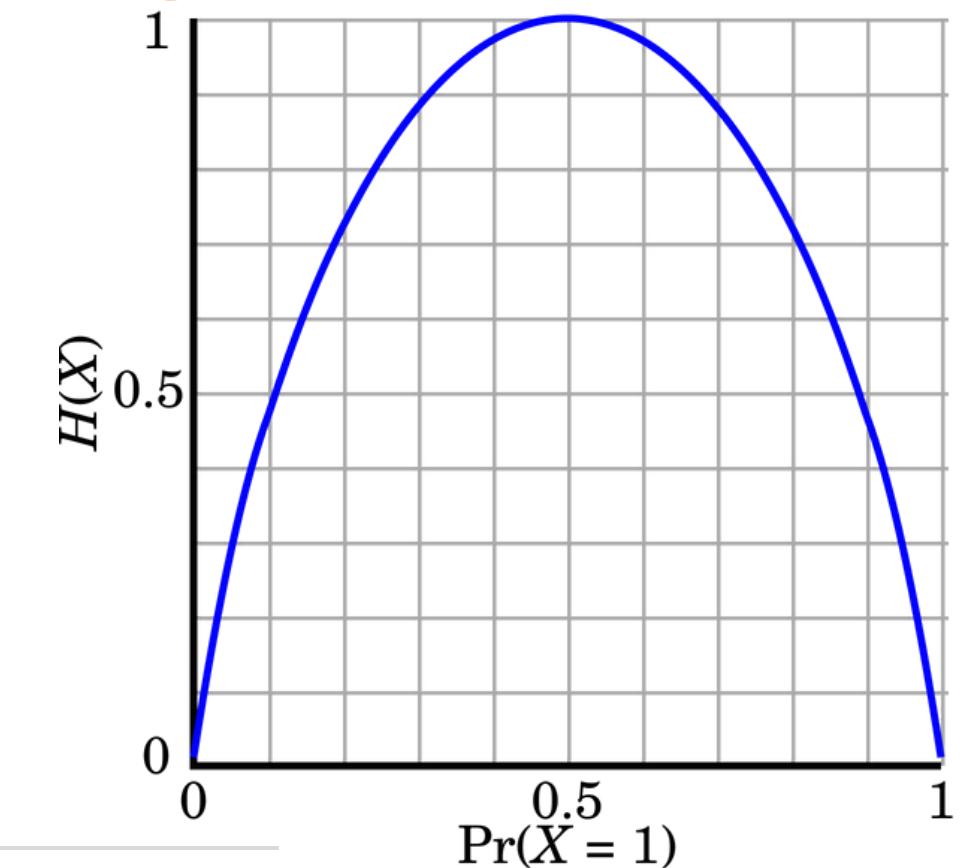
# Decision Trees



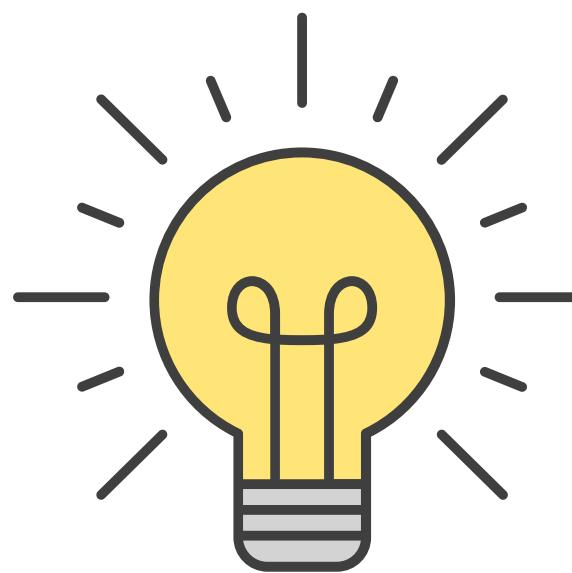
How can we measure how mixed our bags are?

$$\text{Entropy} = -p \log_2 p - (1 - p) \log_2 (1 - p)$$

Basket Contents	Entropy	How Mixed?
100% lemons	0.0	Not mixed – only lemons
50% lemons / 50% oranges	1.0 (max)	Very mixed – max uncertainty
80% lemons / 20% oranges	~0.72	A bit mixed



# Decision Trees



Information Gain = Entropy(before) – Weighted Entropy(after)

$$\text{Weighted Entropy(after)} = \frac{n_1}{n} \cdot \text{Entropy}_1 + \frac{n_2}{n} \cdot \text{Entropy}_2 + \dots$$

Suppose we have:

- 10 fruits: 5 🍋 lemons, 5 🍊 oranges →

$$\text{Entropy(before)} = 1.0$$

Now we split them:

- Basket A: 4 lemons, 1 orange → Entropy = 0.72
- Basket B: 1 lemon, 4 oranges → Entropy = 0.72

$$\text{Weighted Entropy} = 0.5 \cdot 0.72 + 0.5 \cdot 0.72$$

$$\text{Information Gain} = 1.0 - 0.72 = 0.28$$

We reduced uncertainty, so this is a useful split!

# Decision Trees



```
function BuildTree(data, depth):
    if stopping_condition(data, depth):
        return Leaf(predict_class(data))

    best_split = None
    best_score = ∞

    for feature in data.features:
        for threshold in possible_splits(data, feature):
            left, right = split(data, feature, threshold)

            if left is empty or right is empty:
                continue

            score = weighted_impurity(left, right)

            if score < best_score:
                best_score = score
                best_split = (feature, threshold, left, right)

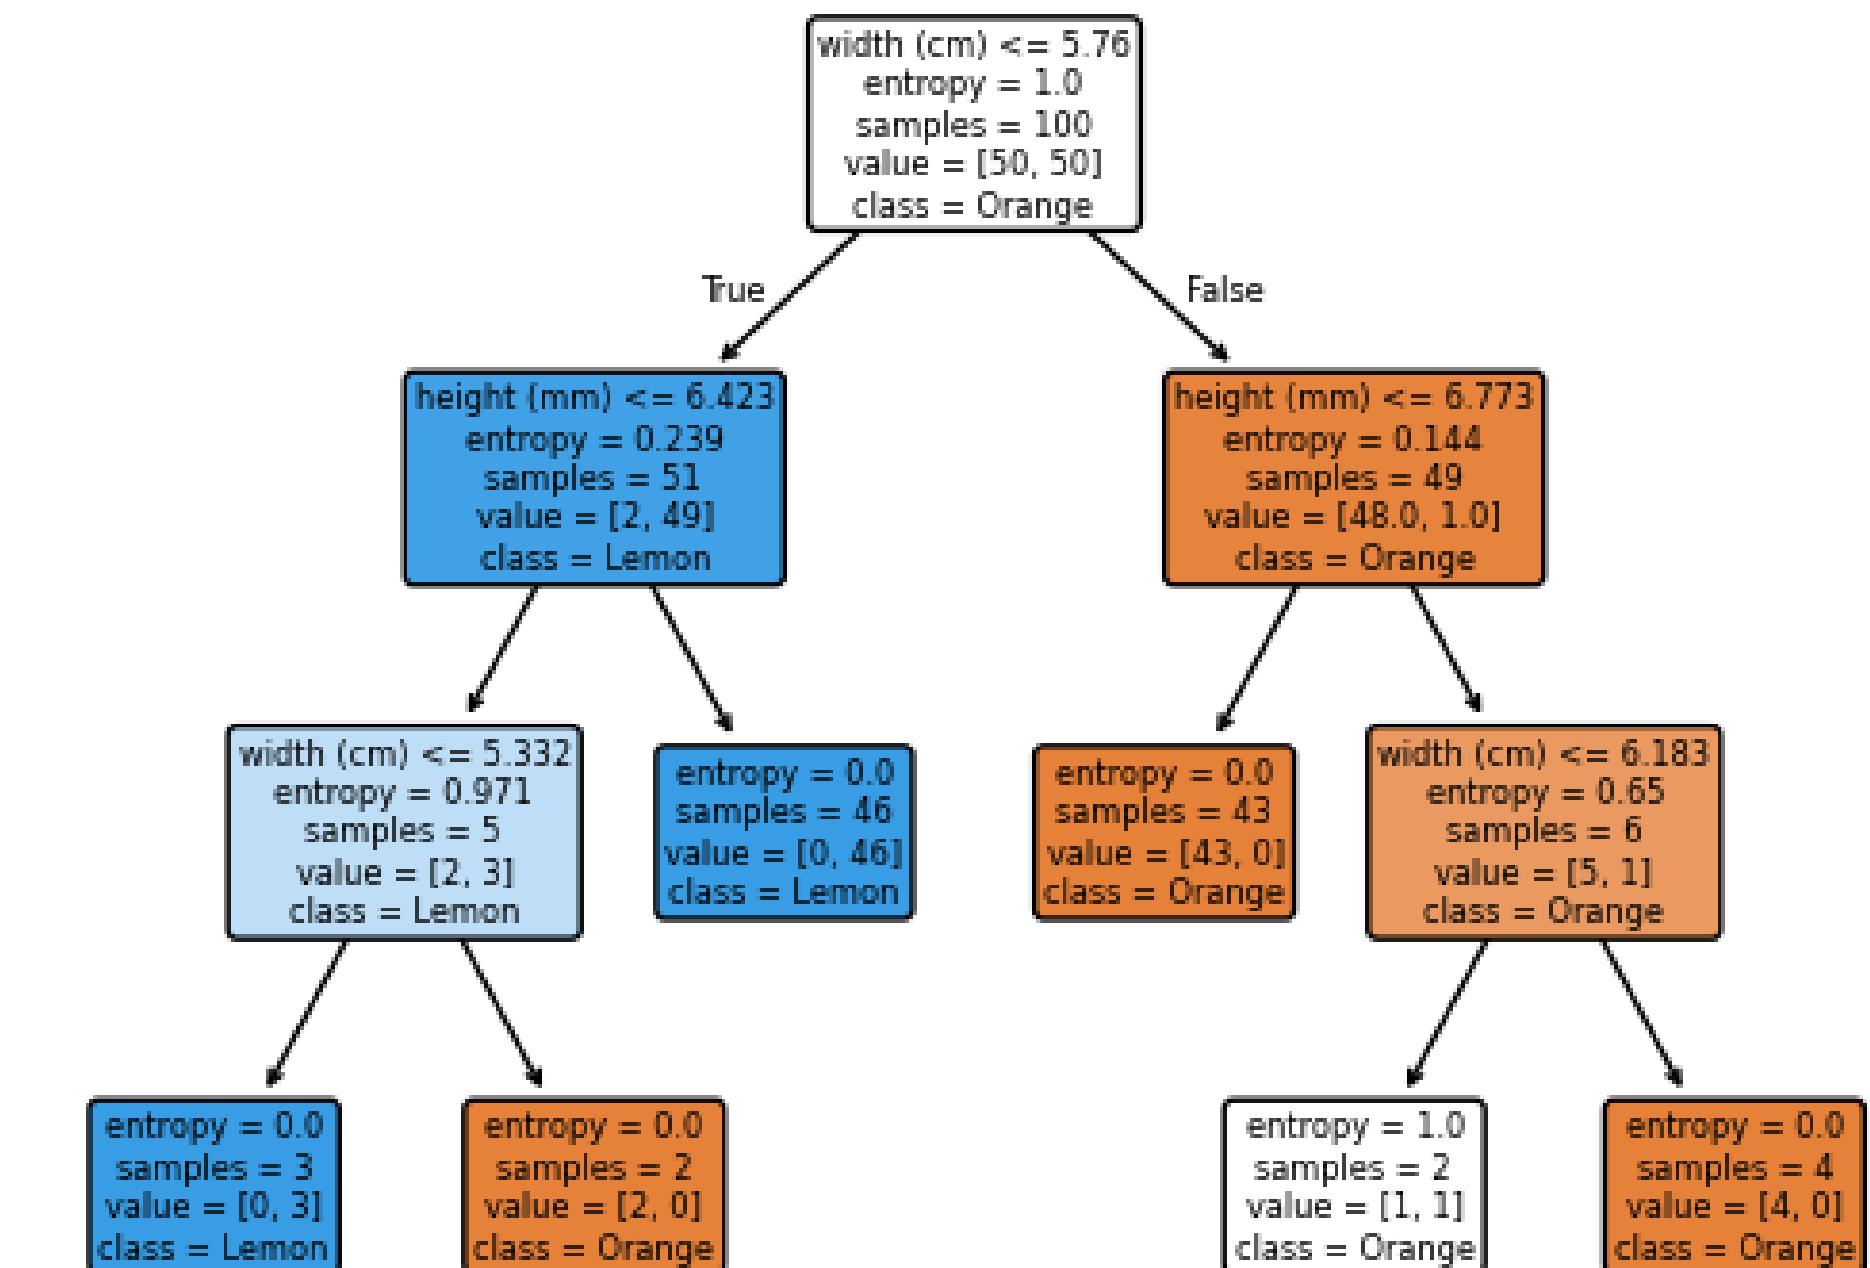
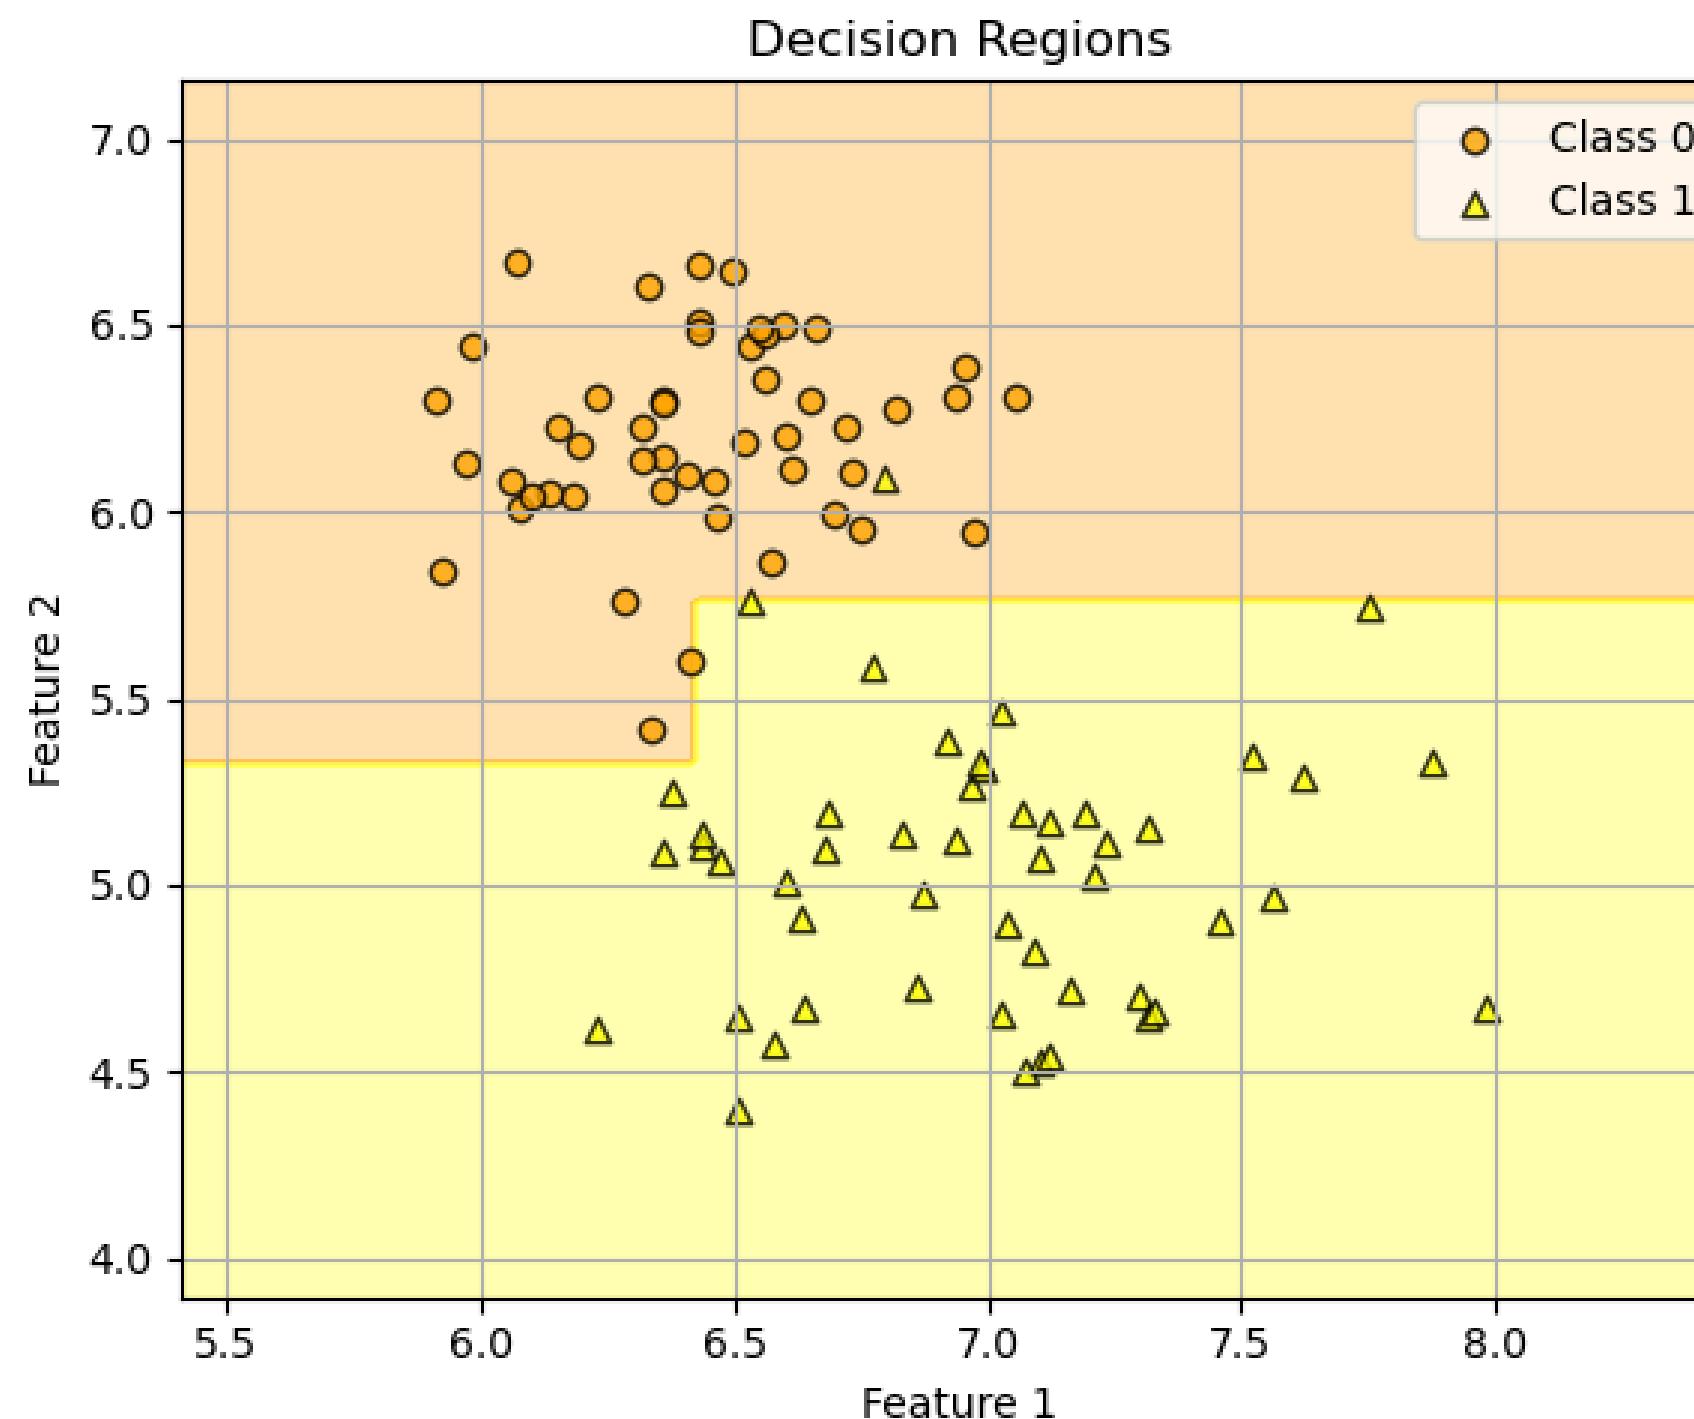
    if best_split is None:
        return Leaf(predict_class(data))

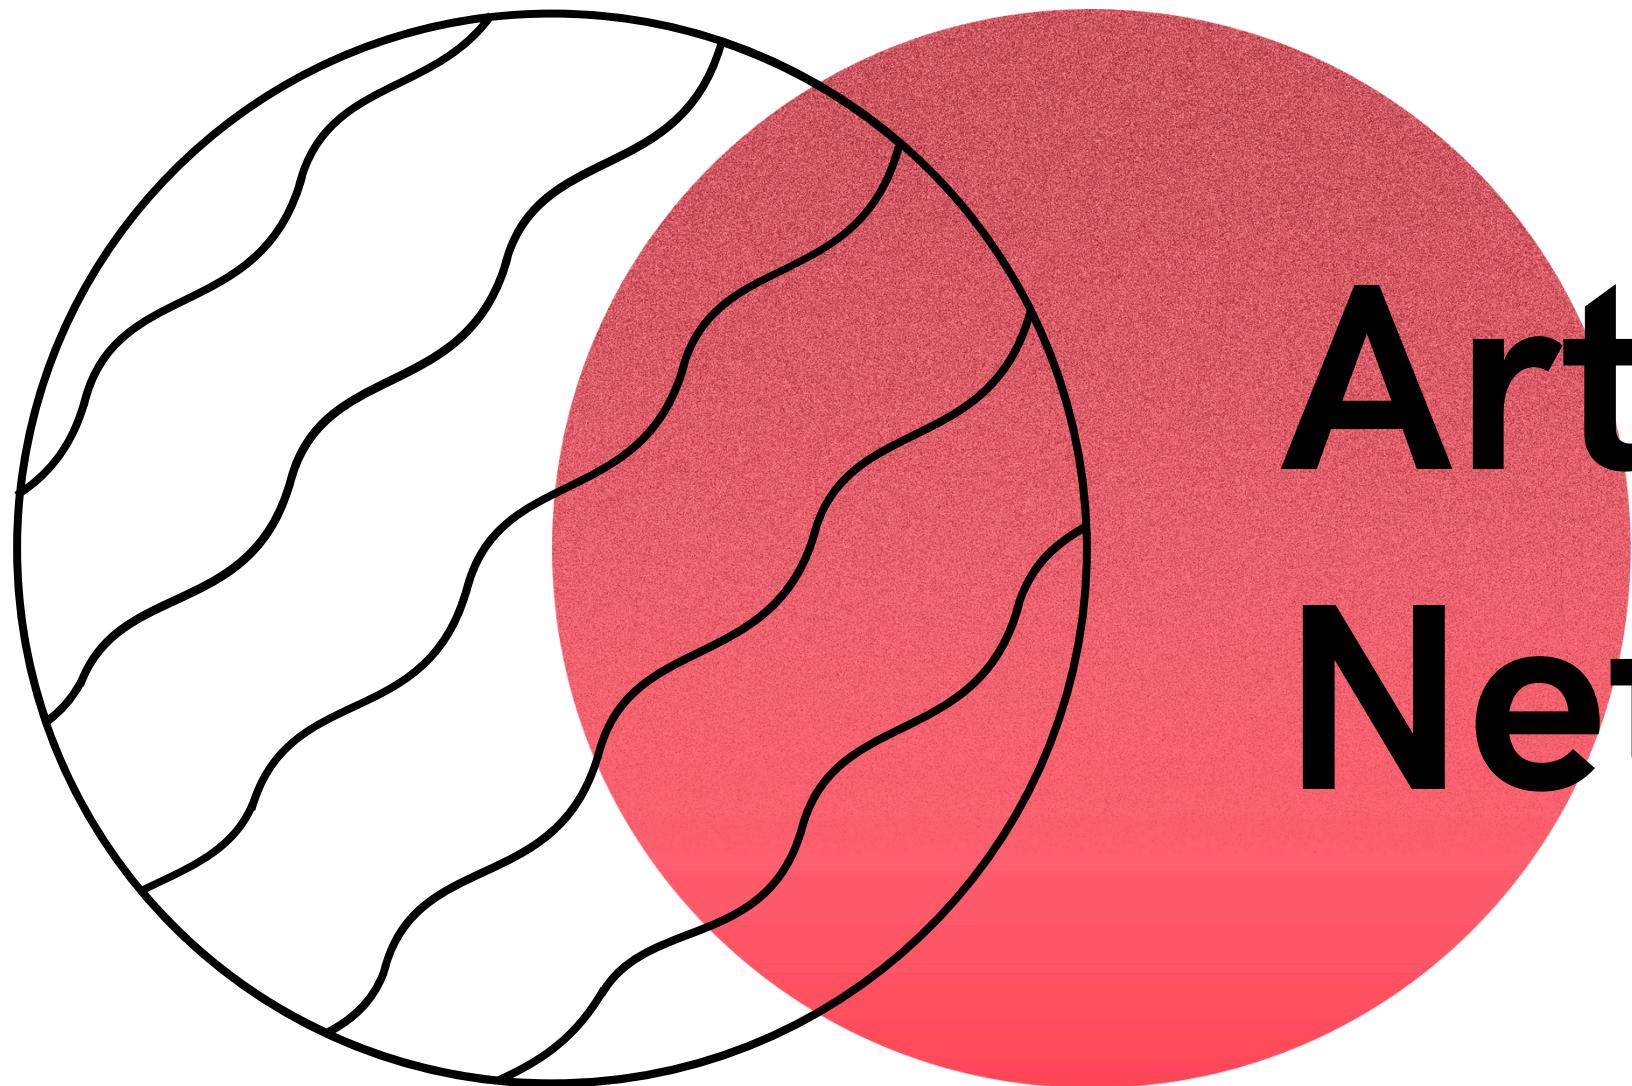
    feature, threshold, left, right = best_split

    left_subtree = BuildTree(left, depth + 1)
    right_subtree = BuildTree(right, depth + 1)

    return Node(feature, threshold, left_subtree, right_subtree)
```

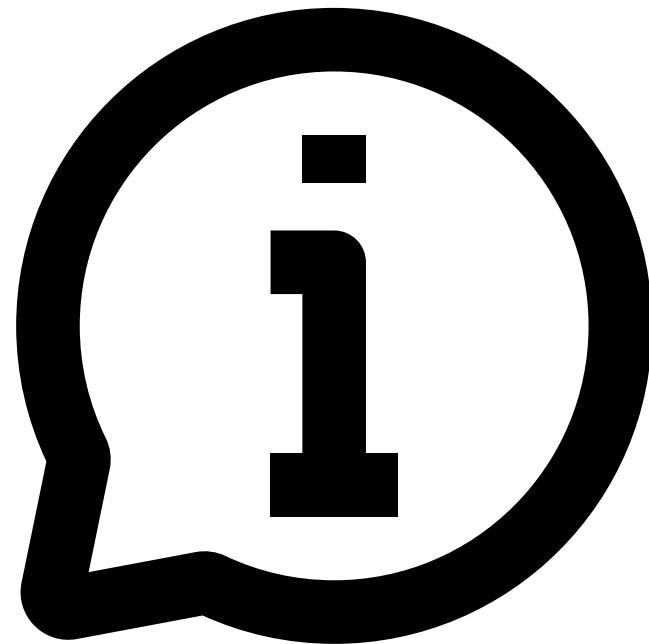
# Decision Trees





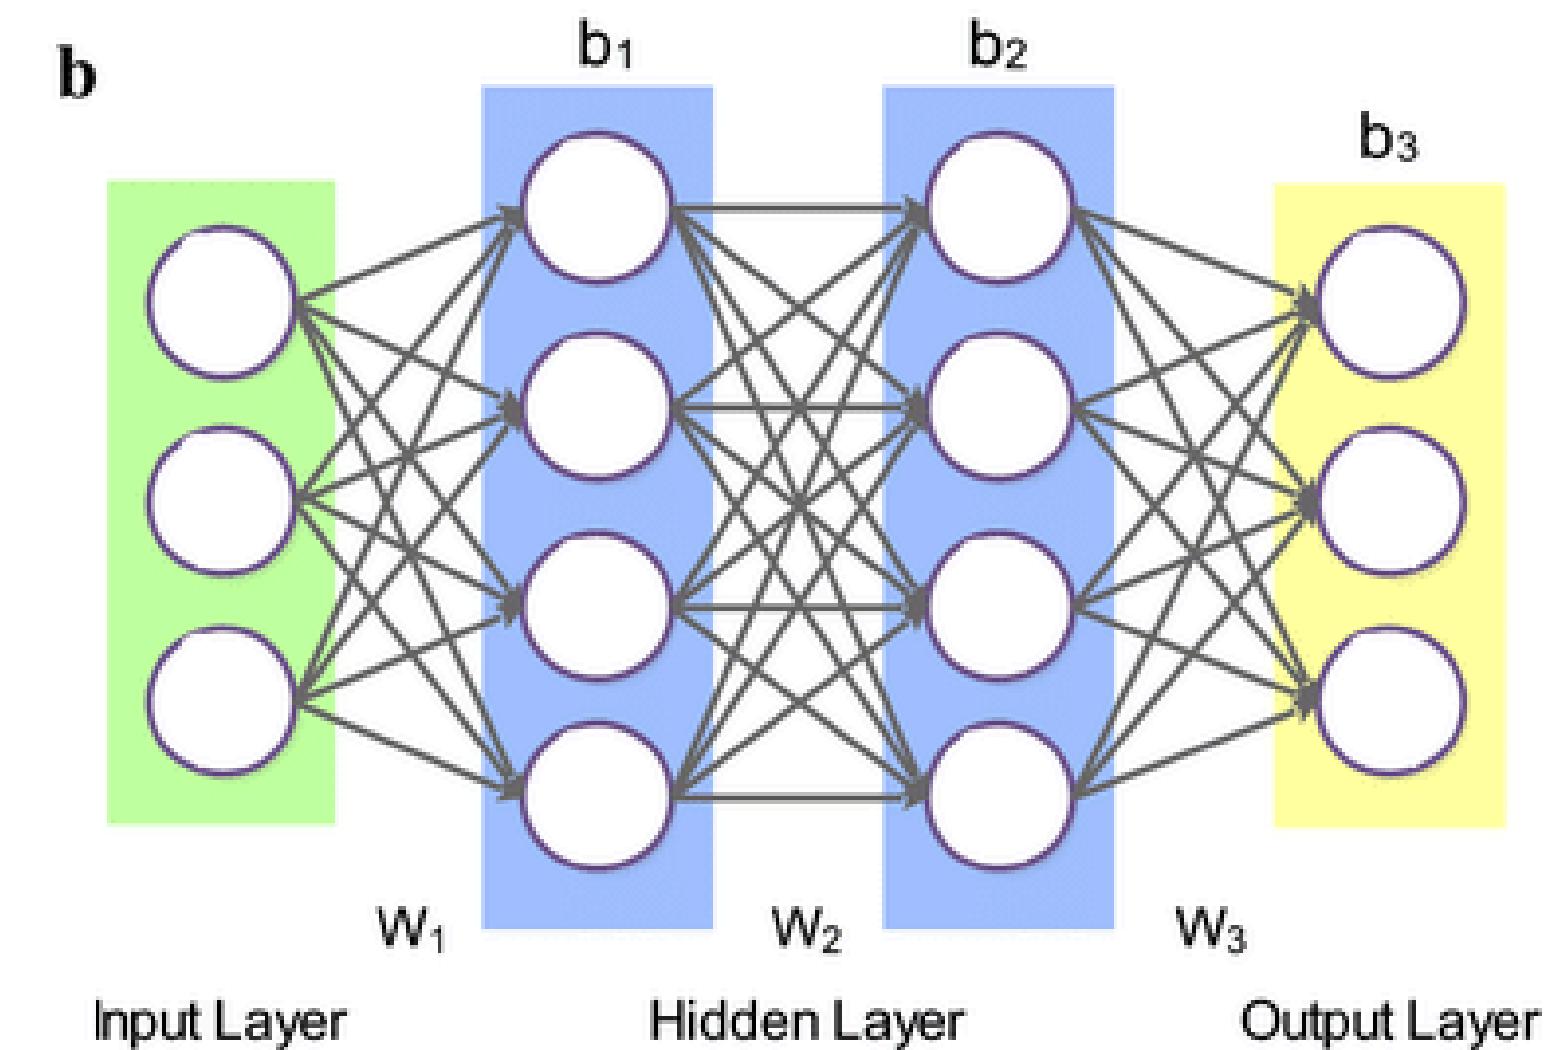
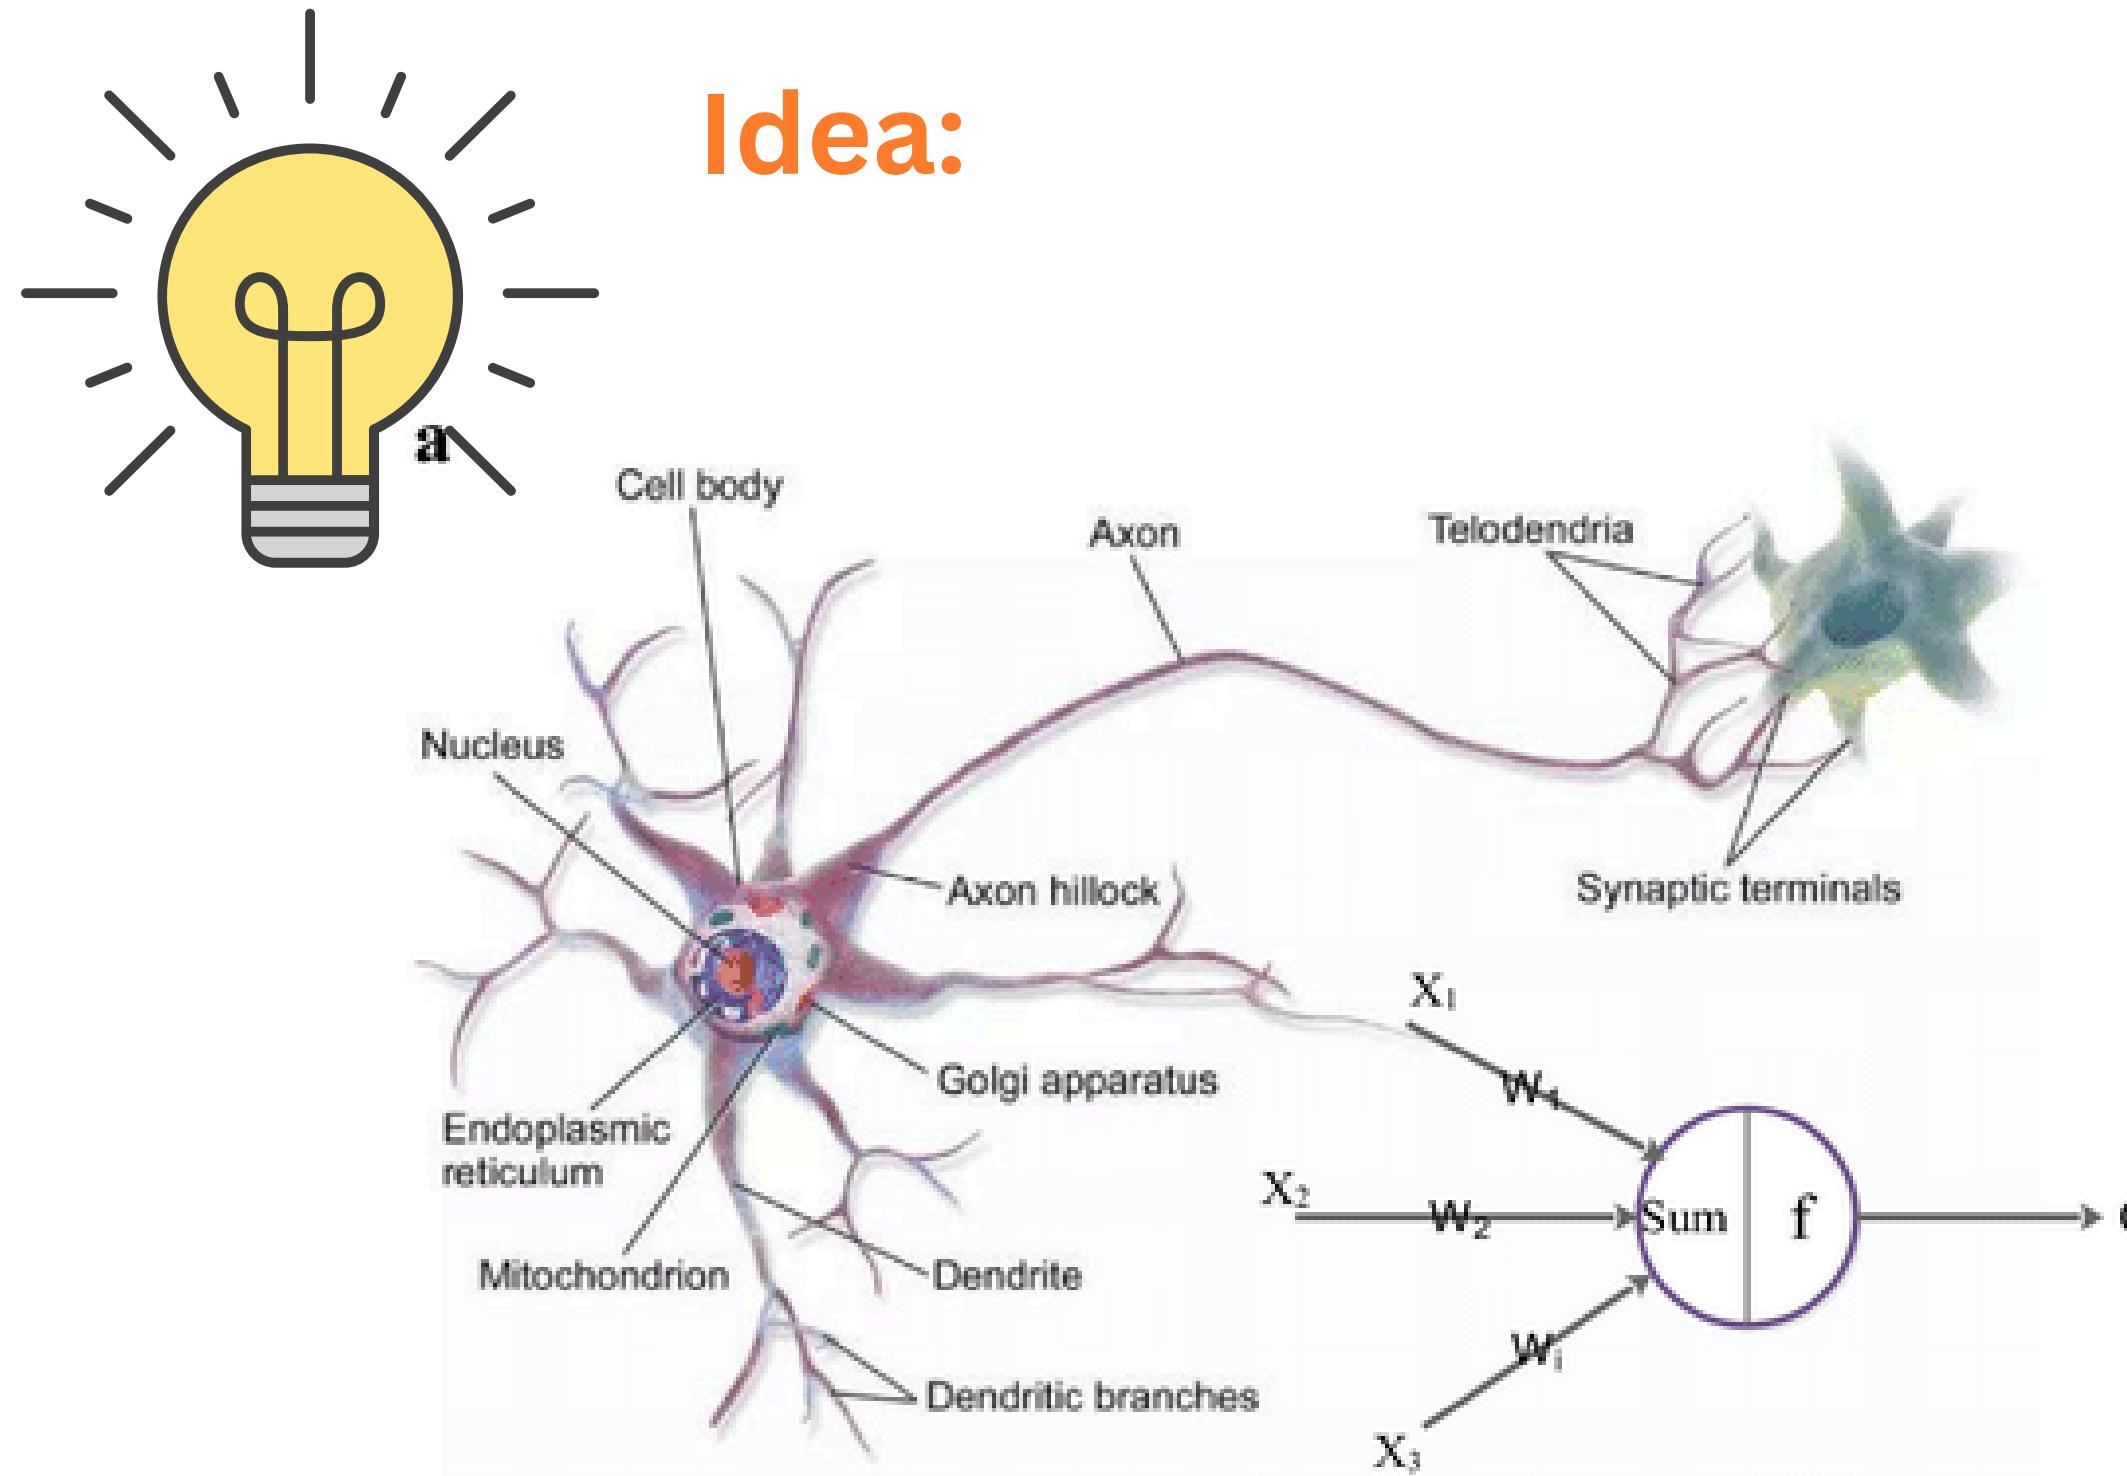
# **Artificial Neural Networks**

# Artificial Neural Networks



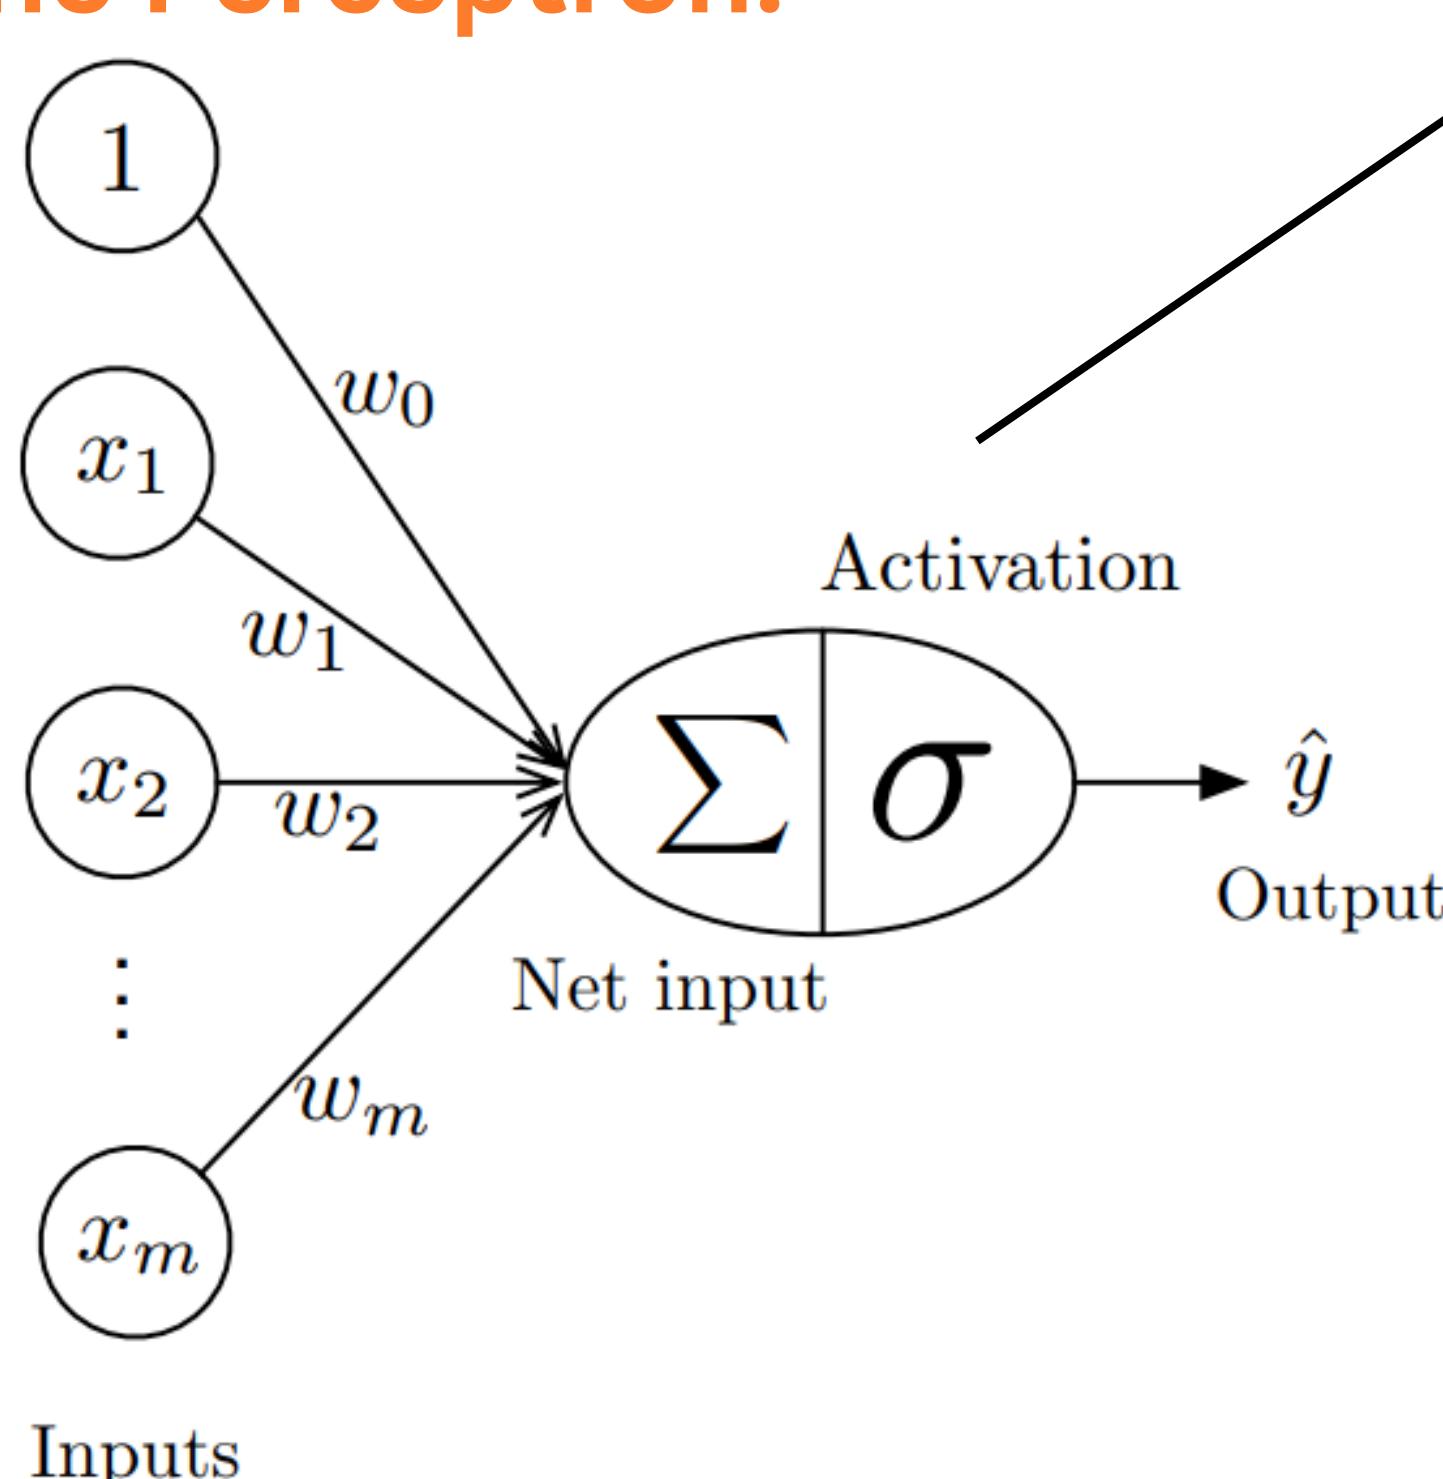
- A computational model ‘inspired’ by how the human brain processes information
- Composed of layers of interconnected “neurons” that transform inputs into outputs
- The ‘Legos’ of Machine Learning

# Artificial Neural Networks



# Artificial Neural Networks

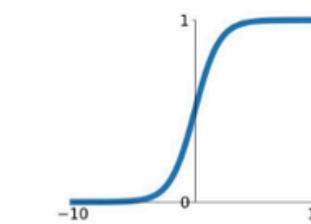
## The Perceptron:



## Activation Functions

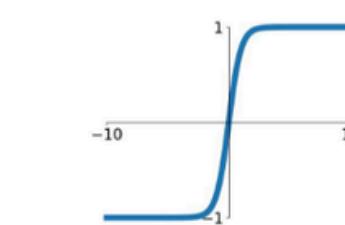
### Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



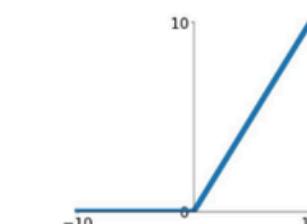
### tanh

$$\tanh(x)$$



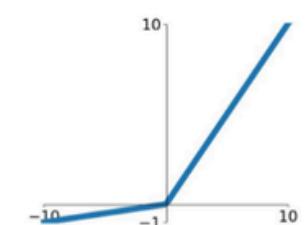
### ReLU

$$\max(0, x)$$



### Leaky ReLU

$$\max(0.1x, x)$$

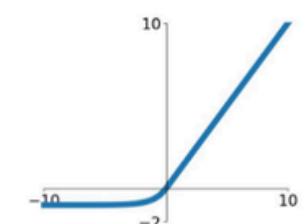


### Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

### ELU

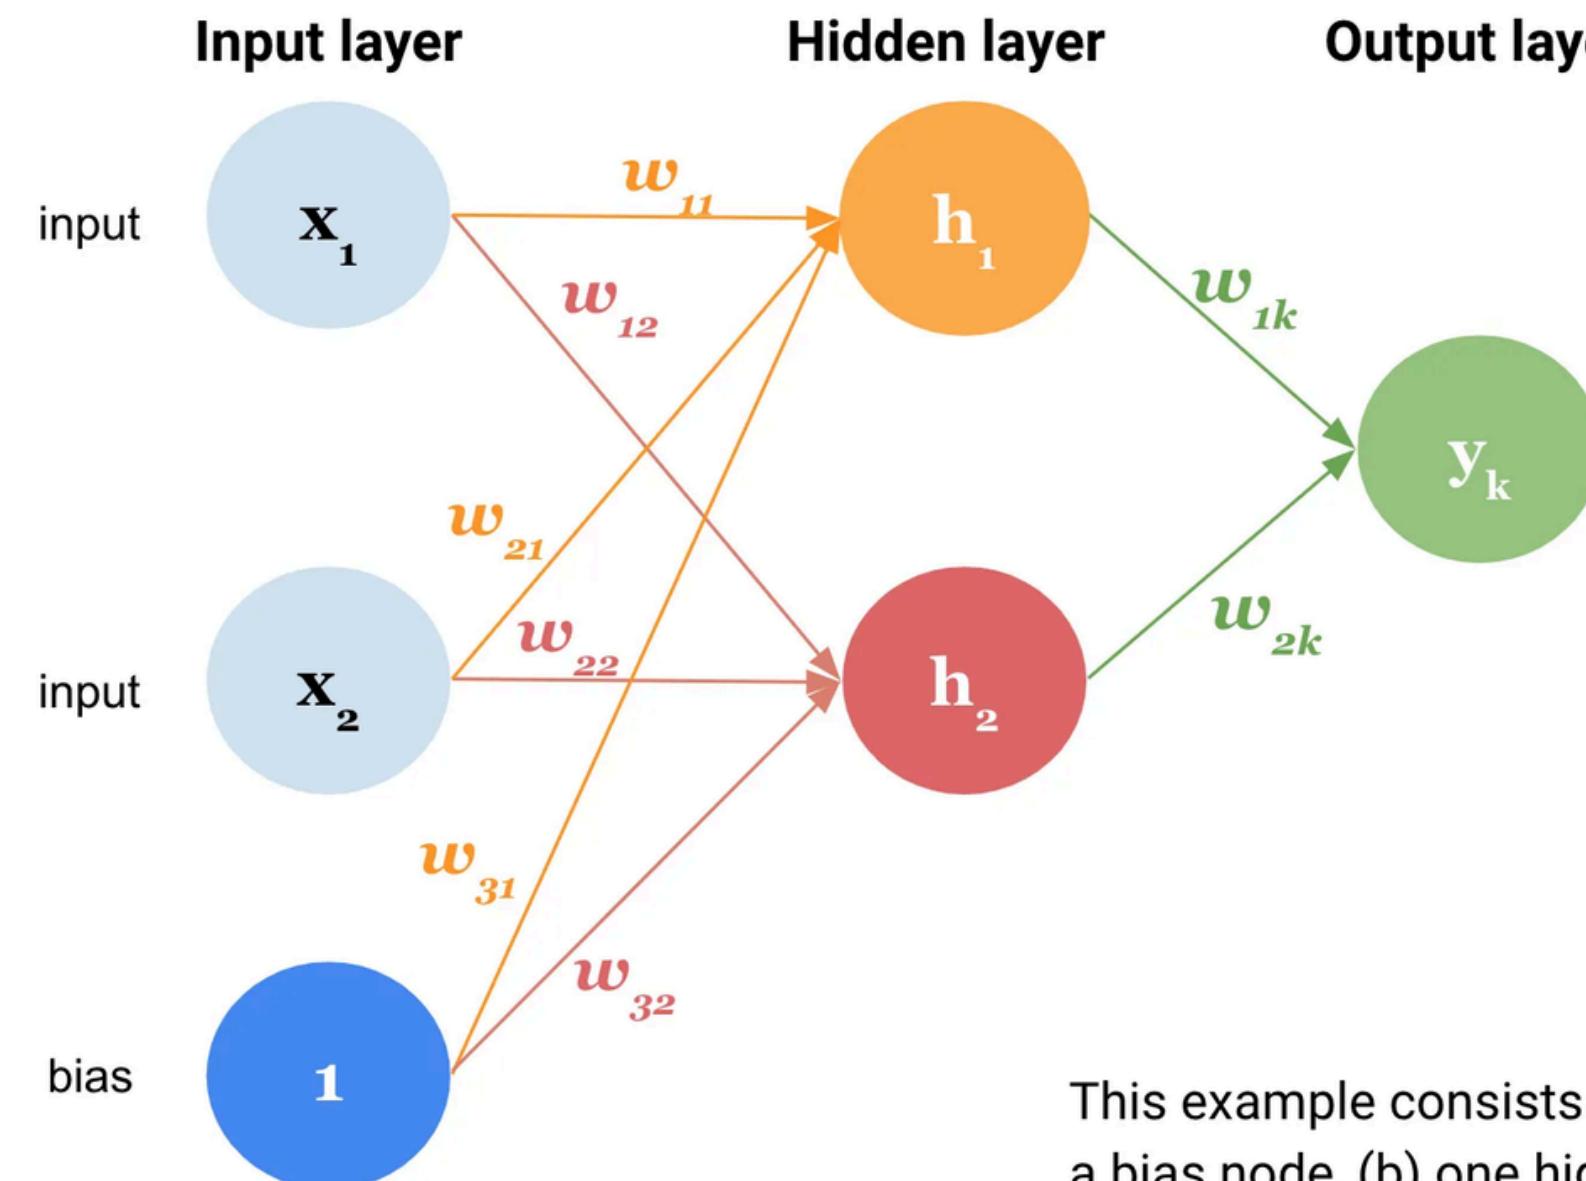
$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



Activation introduces non-linearities

# Artificial Neural Networks

Illustrative example of Multilayer perceptron, a Feedforward neural network



$x_1, x_2$  : input data features  
 $w_{ij}$  : weights of the network  
 $h_1, h_2$  : nodes in the hidden layer  
 $y_k$  : output variable

© AIML.com Research

This example consists of: (a) an input layer with two input nodes and a bias node, (b) one hidden layer with two neurons, and (c) an output layer with one neuron

# Artificial Neural Networks

## Finding the weights:

**1 - Define how bad current predictions are using a Loss Function:**

$$\text{Loss} = -[y \cdot \log(\hat{y}) + (1 - y) \cdot \log(1 - \hat{y})]$$

$y \in \{0, 1\}$  is the **true label**

$\hat{y} \in (0, 1)$  is the **predicted probability**

**2 - Find weights that minimize this Loss:**

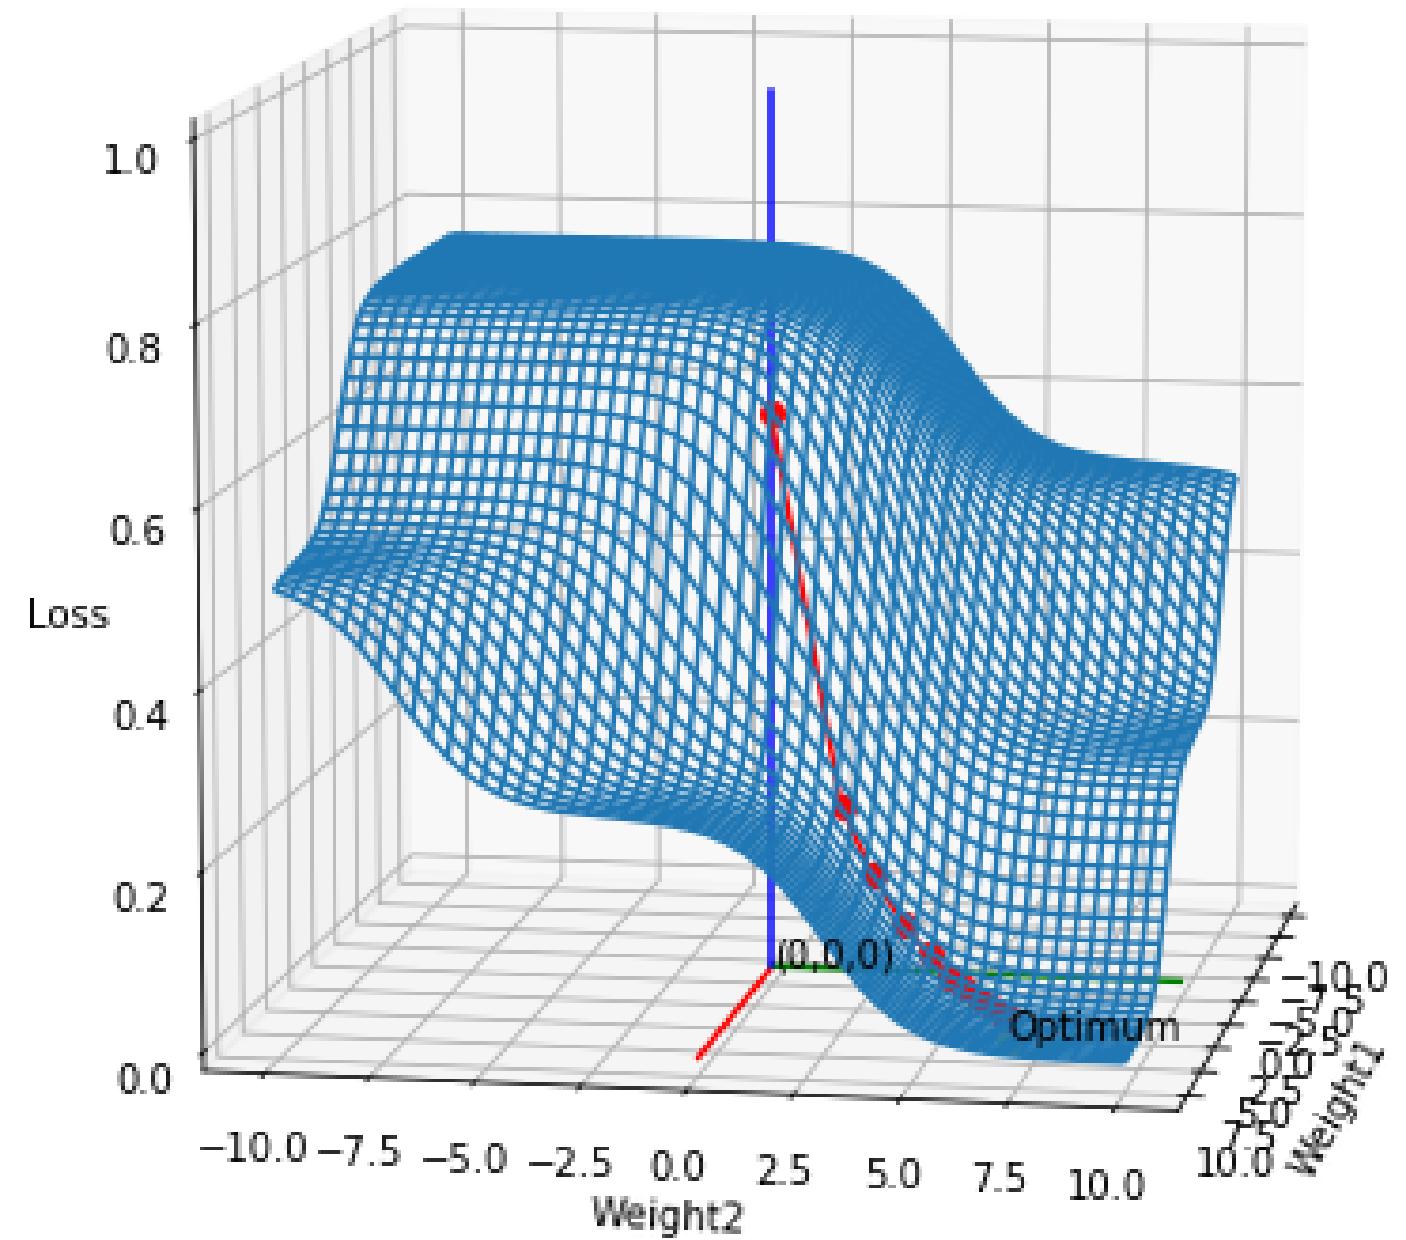
$$\min_{\mathbf{w}} \mathcal{L}(\mathbf{w}) = \frac{1}{n} \sum_{i=1}^n \left[ -y^{(i)} \log(\hat{y}^{(i)}) - (1 - y^{(i)}) \log(1 - \hat{y}^{(i)}) \right]$$



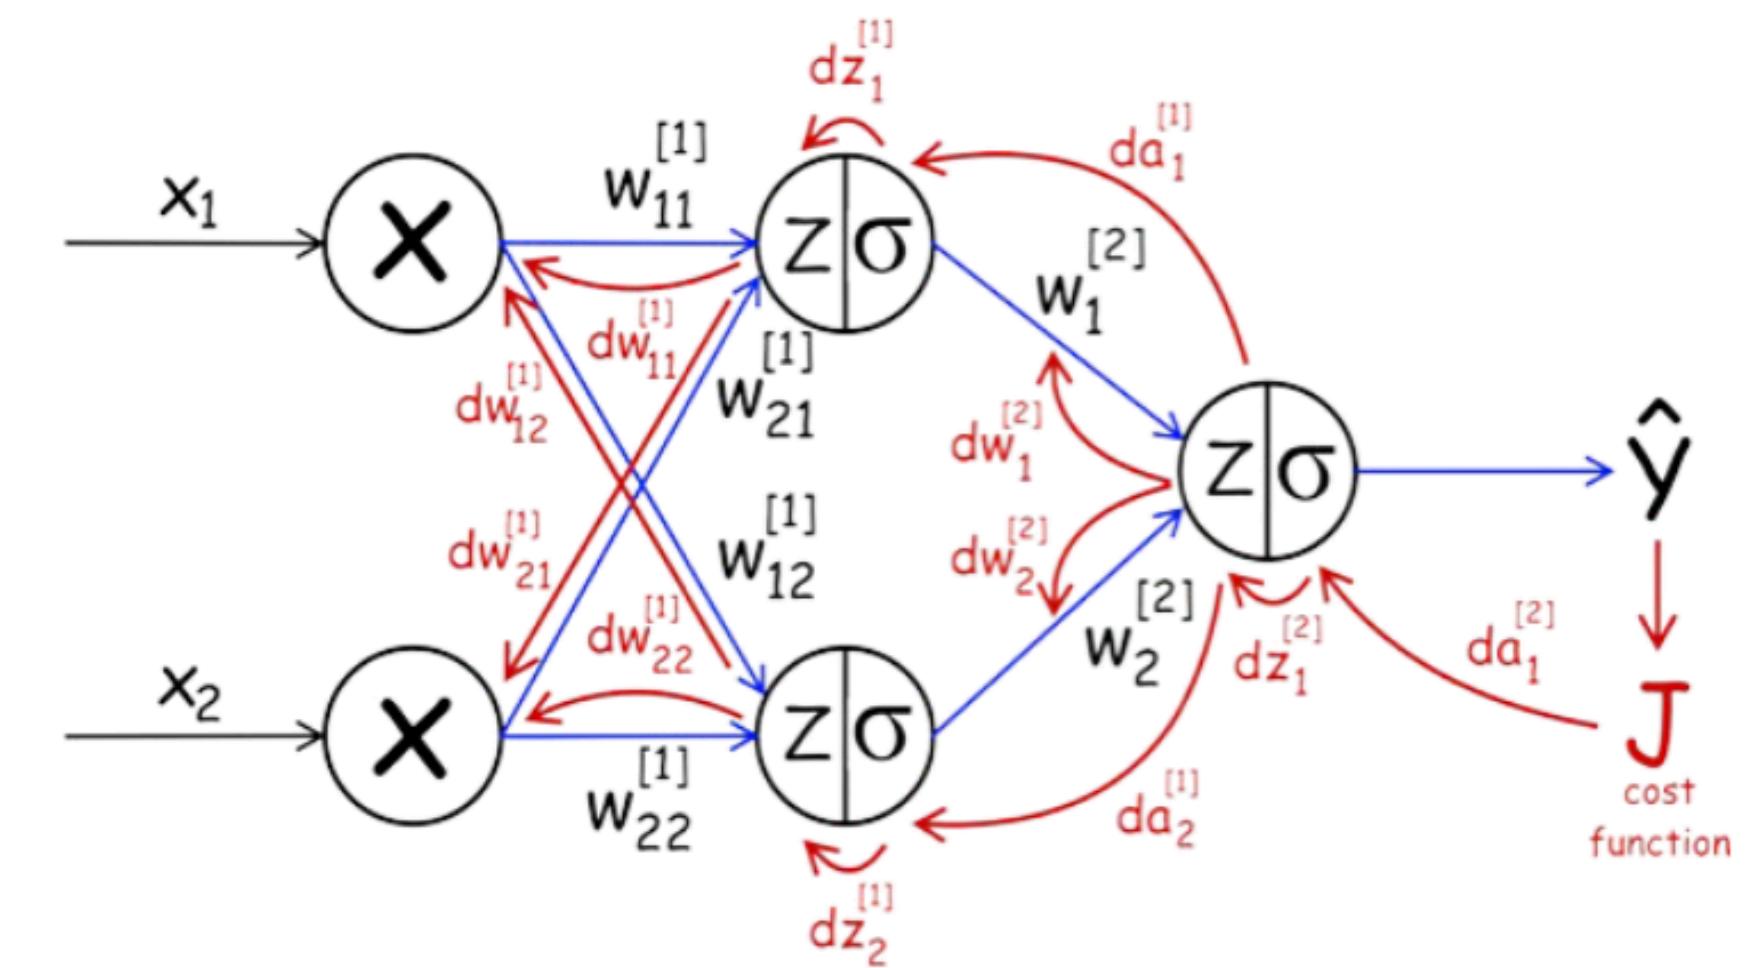
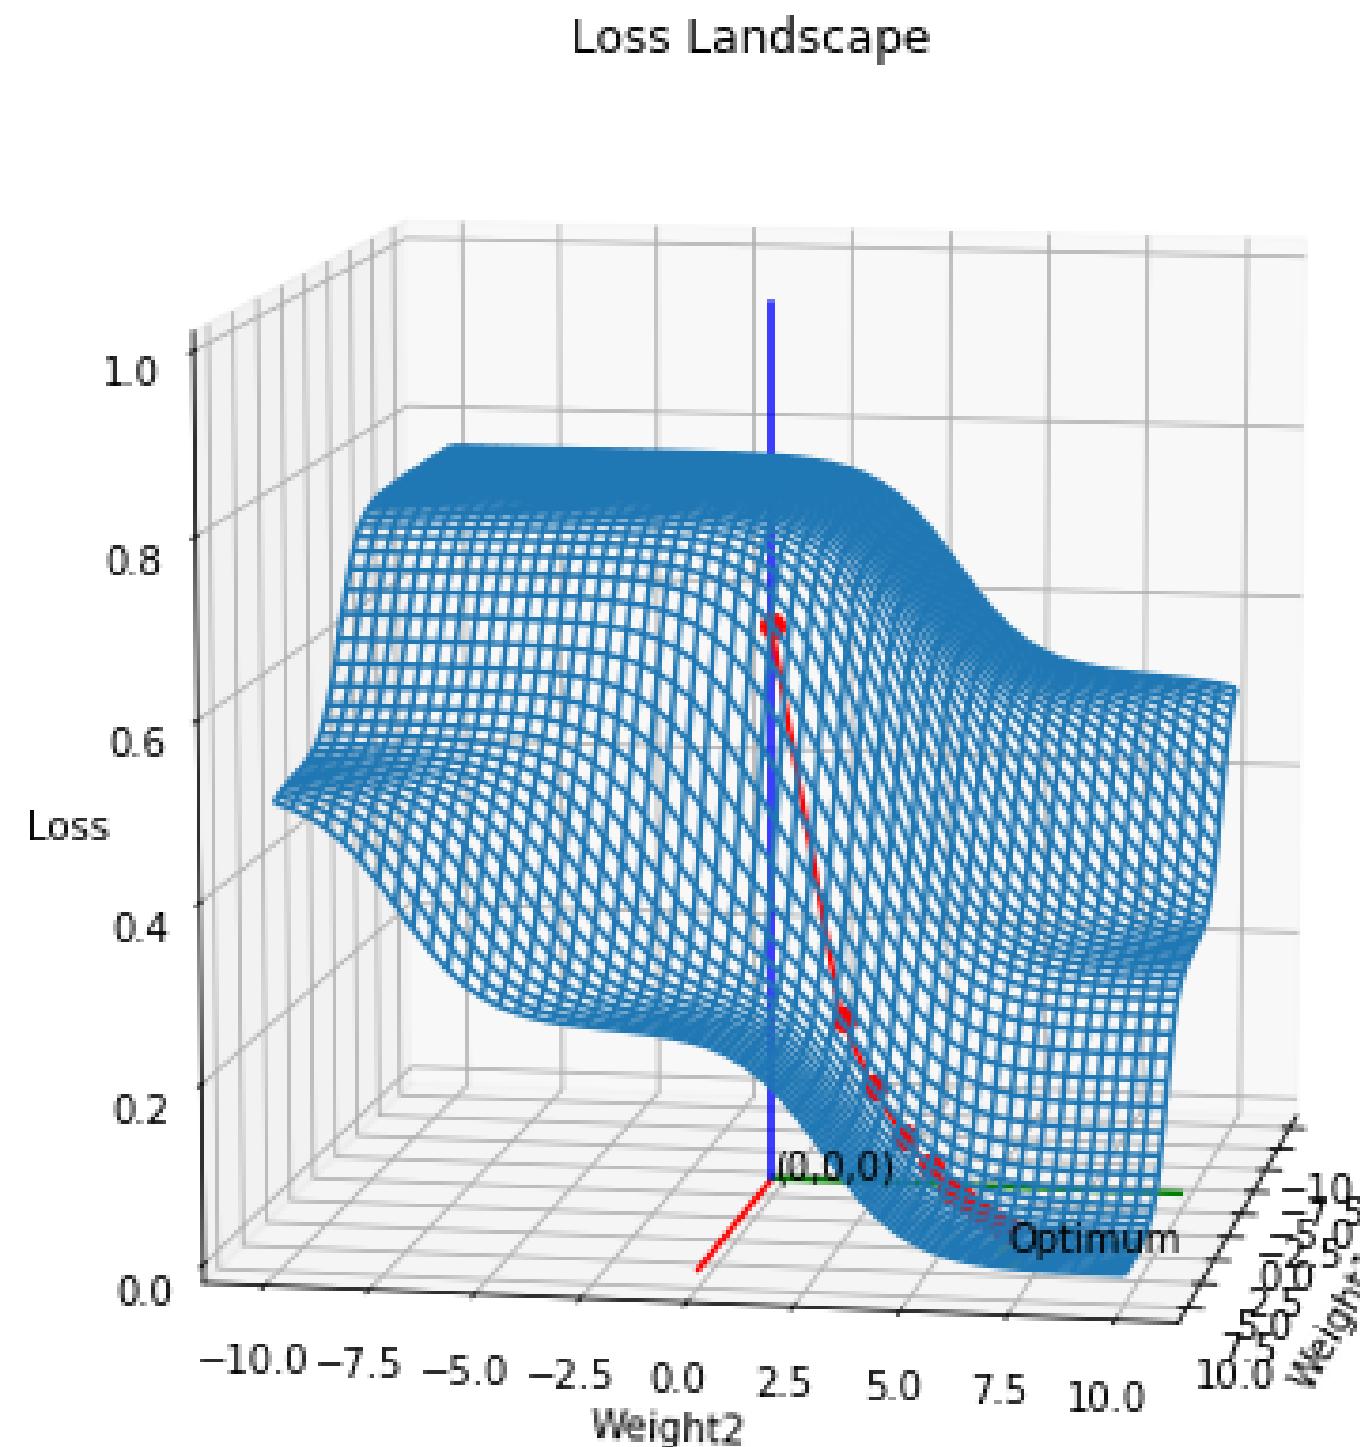
Solved using gradient descent

# Artificial Neural Networks

Loss Landscape



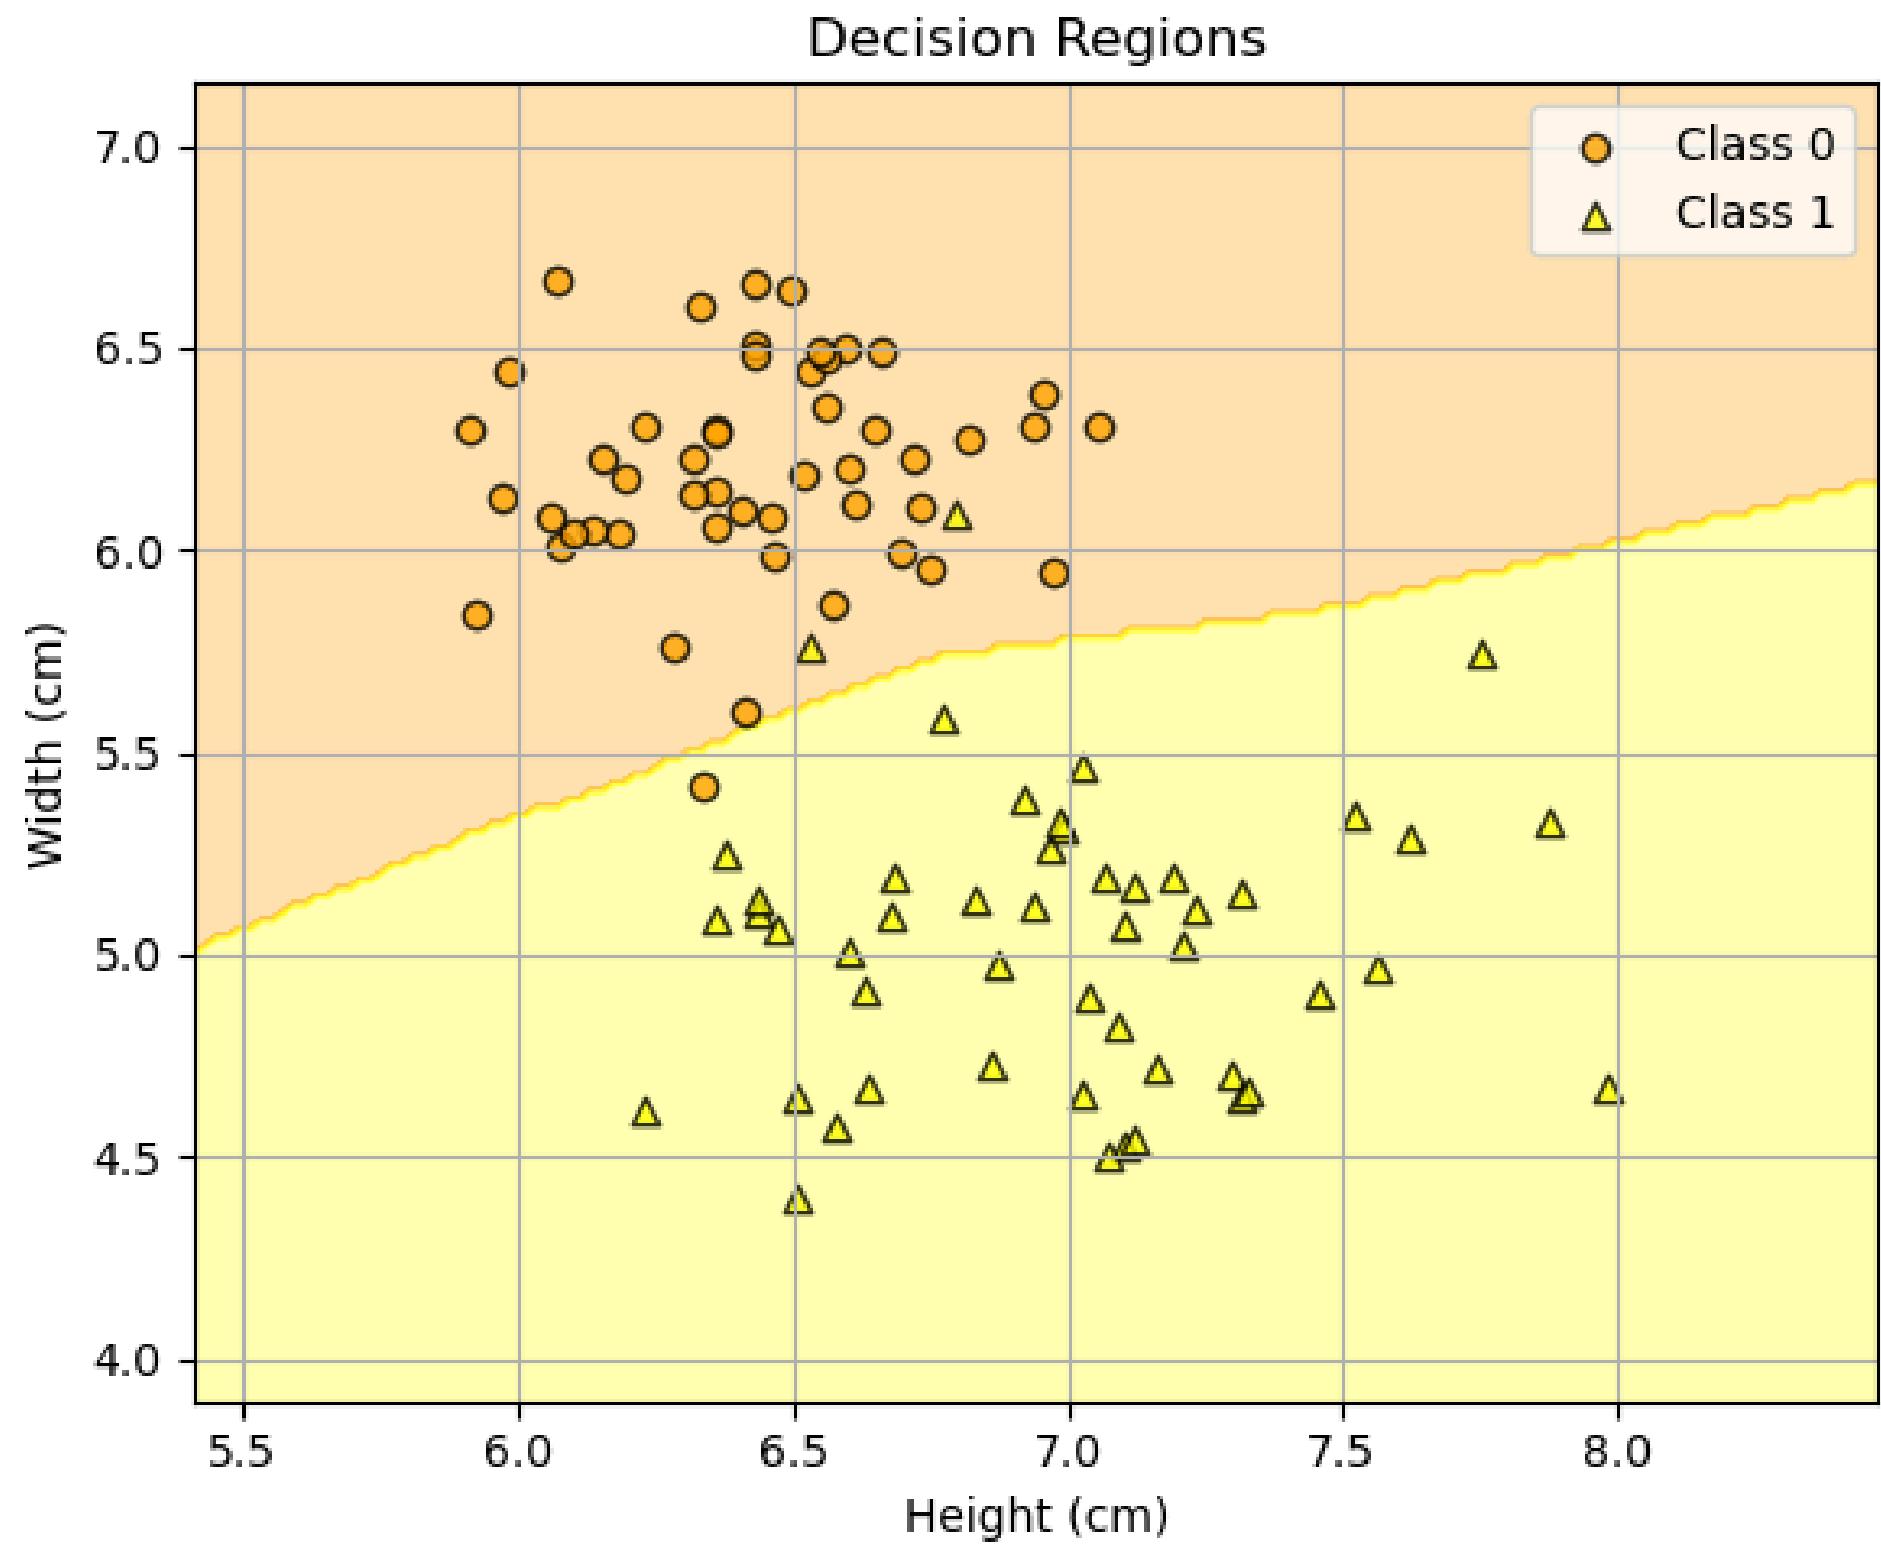
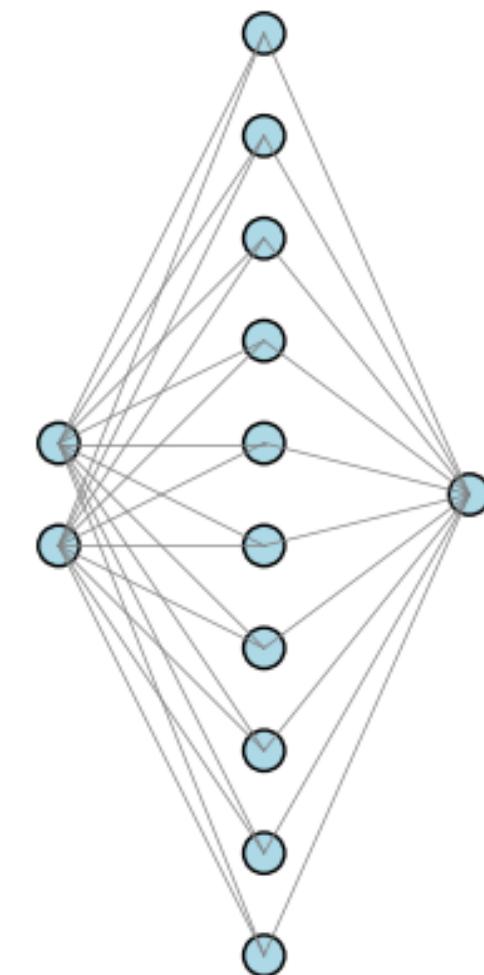
# Artificial Neural Networks

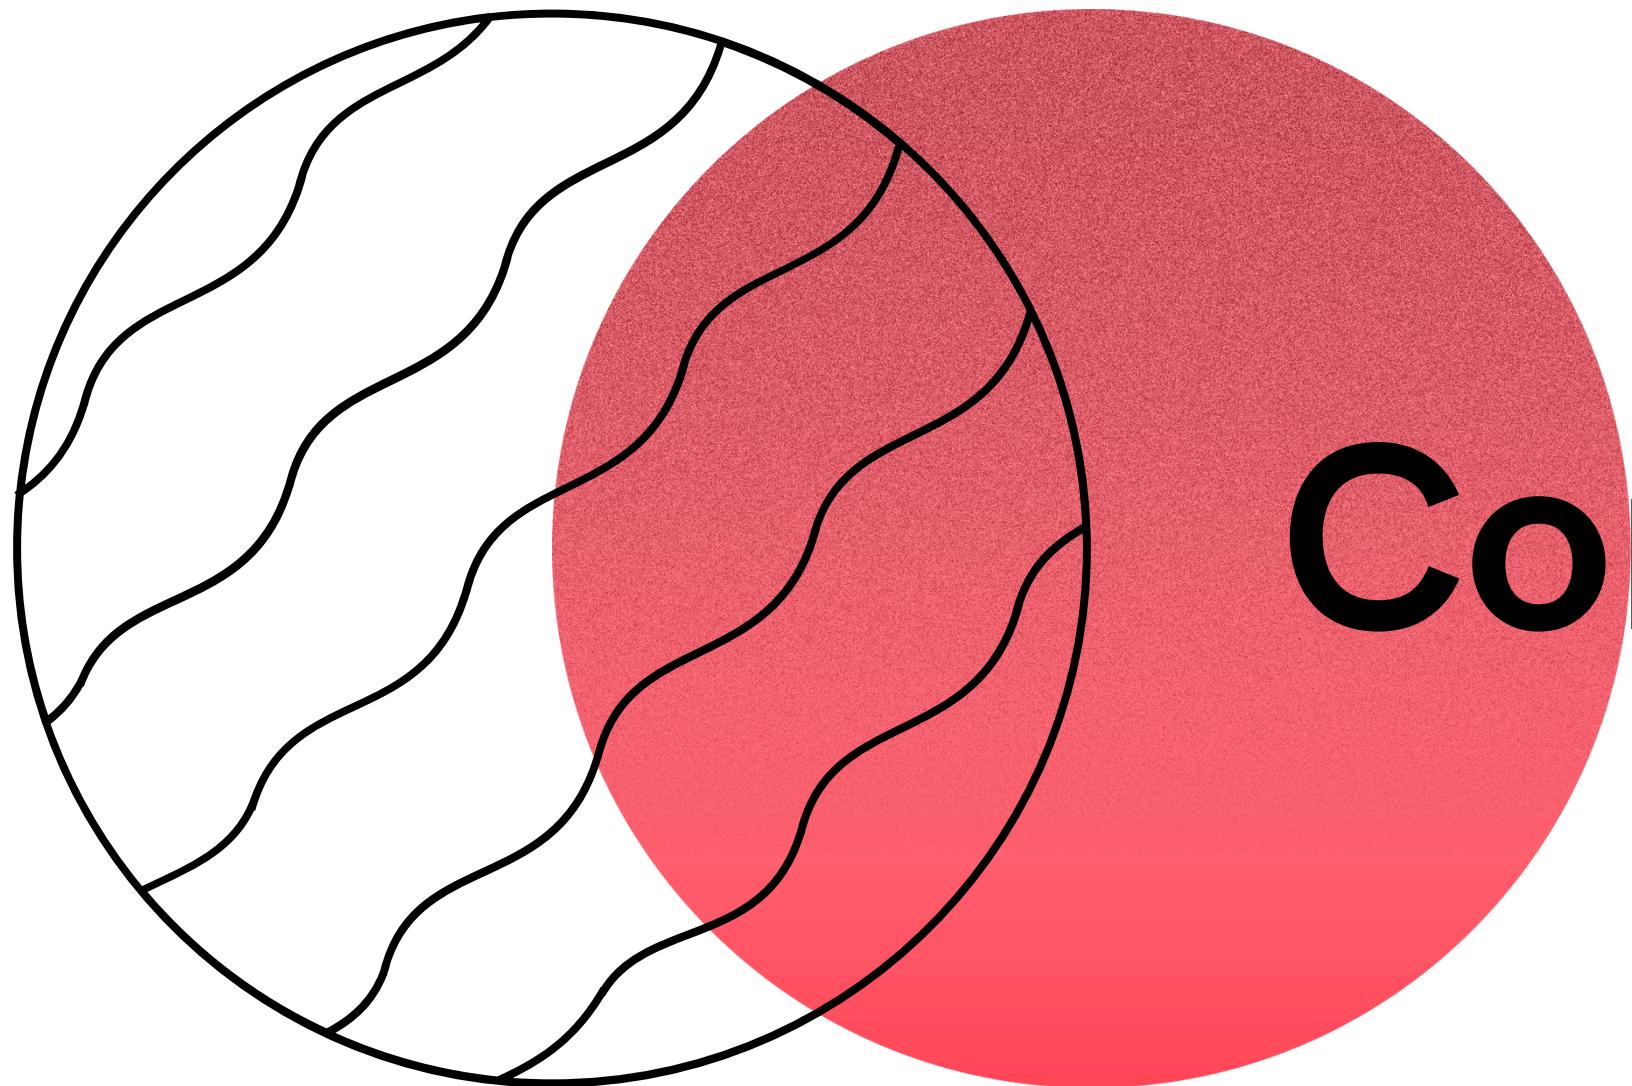


**Backpropagation**

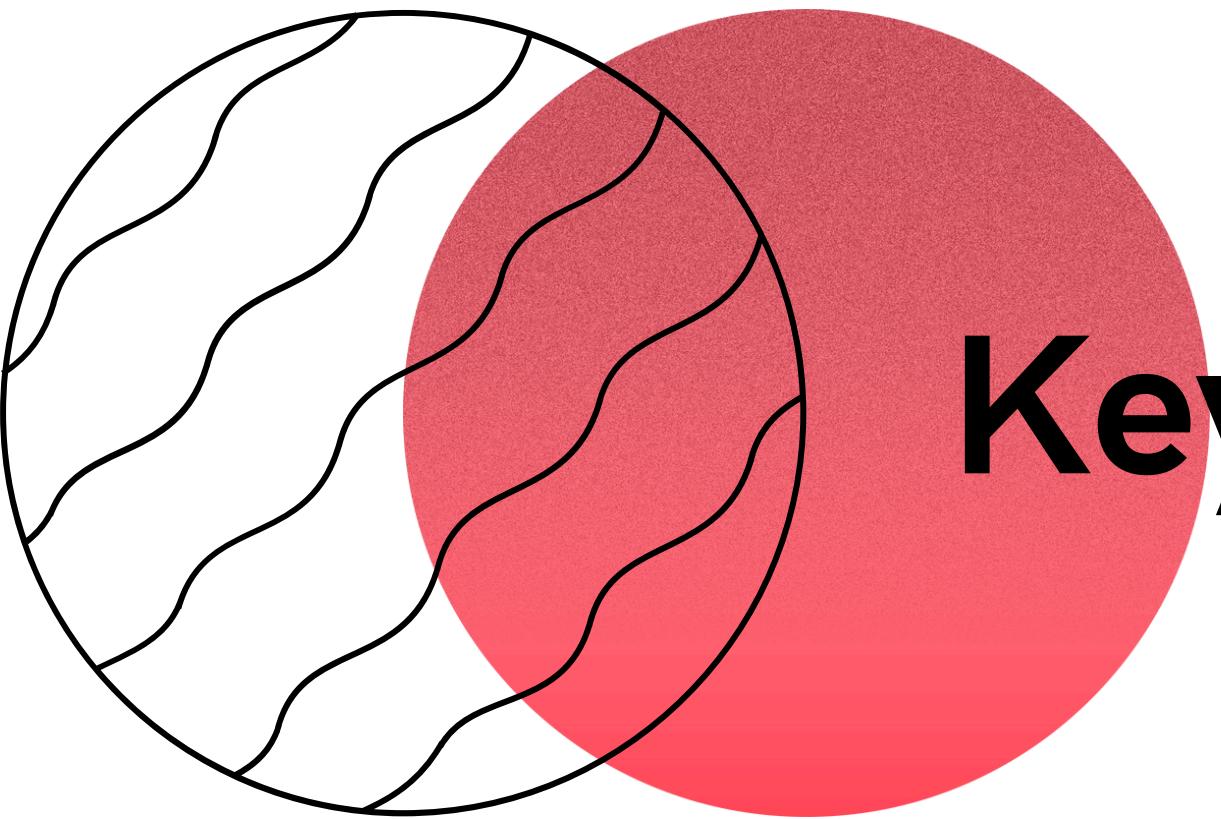
# Artificial Neural Networks

NN Arquitecture





# Conclusion



## Key Takeaways

Feedback here



- K-Nearest Neighbours
- Naive Bayes
- Decision Trees
- Artificial Neural Networks



Evaluating Machine Learning Models